

2. Gene_expression

Fay

2022-05-27

Libraries:

```
library(tidyverse)

## -- Attaching packages ----- tidyverse 1.3.1 --

## v ggplot2 3.3.6      v purrr  0.3.4
## v tibble  3.1.7      v dplyr  1.0.9
## v tidyr   1.2.0      v stringr 1.4.0
## v readr   2.1.2      v forcats 0.5.1

## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()

library(tidyr)
library(dplyr)
library(cowplot)
library(randomForest)

## randomForest 4.7-1.1

## Type rfNews() to see new features/changes/bug fixes.

##
## Attaching package: 'randomForest'

## The following object is masked from 'package:dplyr':
##
##      combine

## The following object is masked from 'package:ggplot2':
##
##      margin

library(ggplot2)
library(caret)
```

```

## Loading required package: lattice

##
## Attaching package: 'caret'

## The following object is masked from 'package:purrr':
##
## lift

library(VIM) # visualizing missing data

## Loading required package: colorspace

## Loading required package: grid

## VIM is ready to use.

## Suggestions and bug-reports can be submitted at: https://github.com/statistikat/VIM/issues

##
## Attaching package: 'VIM'

## The following object is masked from 'package:datasets':
##
## sleep

library(mice) # imputing missing data without predictors

##
## Attaching package: 'mice'

## The following object is masked from 'package:stats':
##
## filter

## The following objects are masked from 'package:base':
##
## cbind, rbind

library(ggpubr)

##
## Attaching package: 'ggpubr'

## The following object is masked from 'package:cowplot':
##
## get_legend

```

```
library(optimx)
```

```
# Here we import the cleaned data set from the previous script derived from the  
# data set challenge infections  
g <- read.csv("output_data/gene_expression/data_products/clean_gene_expression.csv")  
  
# vectors for selecting gene columns  
Genes <- c("IFNy", "CXCR3_bio", "IL.6", "IL.10", "IL.13", "IL.10", "IL.13", "IL1RN",  
           "CASP1", "CXCL9", "IDO1", "IRGM1", "MPO", "MUC2", "MUC5AC", "MYD88",  
           "NCR1", "PRF1", "RETNLB", "SOCS1", "TICAM1", "TNF")
```

Import the data:

```
# we need to change the in challenge infections to a factor  
g$Parasite_challenge <- as.factor(g$Parasite_challenge)  
g$Eim_MC <- as.factor(g$Eim_MC)  
  
# Here I create a new column, where we get the actual infection status  
# According to the melting curve for eimeria  
g <- g %>%  
  dplyr::mutate(current_infection = case_when(  
    Parasite_challenge == "E_ferrisi" & Eim_MC == "TRUE" ~ "E_ferrisi",  
    Parasite_challenge == "E_ferrisi" & Eim_MC == "FALSE" ~ "uninfected",  
    Parasite_challenge == "E_falciformis" & Eim_MC == "TRUE" ~ "E_falciformis",  
    Parasite_challenge == "E_falciformis" & Eim_MC == "FALSE" ~ "uninfected",  
    Parasite_challenge == "uninfected" & Eim_MC == "TRUE" ~ "infected_eimeria",  
    Parasite_challenge == "uninfected" & Eim_MC == "FALSE" ~ "uninfected",  
    TRUE ~ ""  
  ))  
  
# how to impute delta? Replacing with 0 the ones with negative melting curve  
# open for other solutions!  
g <- g %>%  
  dplyr::mutate(Intensity = case_when(  
    Eim_MC == "TRUE" ~ delta,  
    Eim_MC == "FALSE" ~ 0))  
  
# create variable maximum weight loss instead of maximum relative weight loss  
g <- g %>% dplyr::mutate(max_WL = max_WL - 100)
```

Data cleaning

```

#Start by selecting only the genes and the maximum weight loss for each mouse
# Apparently the relative end weight doesn't work so well for predictions

g.1 <- g %>%
  dplyr::select(c(max_WL, all_of(Genes)))

# to get reproducible results we use a seed
set.seed(42)

# We want the maximum weight loss to be predicted by the data in all of the other columns

# iter = how many random forests are needed, in theory 6 are enough
g.imputed <- rfImpute(max_WL ~ ., data = g.1, iter = 6)

```

Imputing missing data + cleaning

```

##      |      Out-of-bag |
## Tree |      MSE %Var(y) |
## 300 |      26.1  61.11 |
##      |      Out-of-bag |
## Tree |      MSE %Var(y) |
## 300 |      27.08  63.39 |
##      |      Out-of-bag |
## Tree |      MSE %Var(y) |
## 300 |      27.97  65.49 |
##      |      Out-of-bag |
## Tree |      MSE %Var(y) |
## 300 |      28.27  66.19 |
##      |      Out-of-bag |
## Tree |      MSE %Var(y) |
## 300 |      28.26  66.17 |
##      |      Out-of-bag |
## Tree |      MSE %Var(y) |
## 300 |      28.24  66.11 |

```

```

g.imputed <- g.imputed %>% dplyr::select(-max_WL)

g_minus <- g %>%
  dplyr::select(-all_of(Genes))

#full data set containing the imputed gene expression data
g.imputed <- cbind(g_minus, g.imputed)

```

How many mice are in the infection planning?

```

g.imputed %>%
  filter(infection == "challenge") %>%
  group_by(Parasite_challenge) %>%
  summarize(length(EH_ID))

```

```
## # A tibble: 3 x 2
```

```
##   Parasite_challenge 'length(EH_ID)'
##   <fct>                <int>
## 1 E_falciformis        22
## 2 E_ferrisi            47
## 3 uninfected           47
```

How many mice are indeed infected?

```
g.imputed %>%
  filter(infection == "challenge") %>%
  group_by(current_infection) %>%
  summarize(length(EH_ID))
```

```
## # A tibble: 4 x 2
##   current_infection 'length(EH_ID)'
##   <chr>                <int>
## 1 E_falciformis        22
## 2 E_ferrisi            39
## 3 infected_eimeria     9
## 4 uninfected          46
```

I guess mice got mixed up here?

Splitting data into training and testing sets Splitting between training and testing: - Assess model performance on unseen data - Avoid over-fitting

```
#select the relevant columns:
g.imputed <- g.imputed %>%
  dplyr::select(c(max_WL, all_of(Genes))) %>%
  dplyr::select(- CXCR3_bio)

# split data into training and test

set.seed(123) # this will help us reproduce this random assignment

# in this way we can pick the random numbers

training.samples <- g.imputed$max_WL%>%
  createDataPartition(p = .7, # this is the partiicipation! In this case 0.7 = training data and 0.3 = te
    list = FALSE) # we don't want to get a list in return

train.data <- g.imputed[training.samples, ] #include all the randomly selected rows
test.data <- g.imputed[-training.samples, ]
```

```
#train the model
model <- randomForest(max_WL ~., data = train.data, proximity = TRUE,
  ntree = 1000) # number of trees


print(model)
```

Building the model

```
##  
## Call:  
## randomForest(formula = max_WL ~ ., data = train.data, proximity = TRUE,      ntree = 1000)  
##           Type of random forest: regression  
##           Number of trees: 1000  
## No. of variables tried at each split: 6  
##  
##           Mean of squared residuals: 24.8344  
##           % Var explained: 40.55
```

Plotting the model will illustrate the error rate as we average across more trees and shows that our error rate stabilizes with around 200 trees.

```
plot(model)
```



The plotted error rate above is based on the OOB sample error and can be accessed directly at `m1$mse`. Thus, we can find which number of trees providing the lowest error rate, which is 257 trees providing an weight error of 5.024738.

```
# number of trees with lowest MSE  
which.min(model$mse)
```

```
## [1] 71
```

```
## [1] 257
```

```
# RMSE of this optimal random forest  
sqrt(model$mse[which.min(model$mse)])
```

```
## [1] 4.88103
```

```
## [1] 5.024738
```

<https://uc-r.github.io/s>

RandomForest also allows us to use a validation set to measure predictive accuracy if we did not want to use the OOB samples.

Tutorial: <https://hackernoon.com/random-forest-regression-in-r-code-and-interpretation> Random forest regression in R provides two outputs: decrease in mean square error (MSE) and node purity. Prediction error described as MSE is based on permuting out-of-bag sections of the data per individual tree and predictor, and the errors are then averaged. In the regression context, Node purity is the total decrease in residual sum of squares when splitting on a variable averaged over all trees (i.e. how well a predictor decreases variance). MSE is a more reliable measure of variable importance. If the two importance metrics show different results, listen to MSE. If all of your predictors are numerical, then it shouldn't be too much of an issue

Mean Decrease Gini (IncNodePurity) - This is a measure of variable importance based on the Gini impurity index used for the calculating the splits in trees.

Improving Your Model Your model depends on the quality of your dataset and the type of Machine Learning algorithm used. Therefore, to improve the accuracy of your model, you should:

Check what attributes affect our model the most and what variables to leave out in future analysis Find out what other attributes affect a person's wage; we can use as predictors in future analysis Tweak the algorithm (e.g. change the ntree value) Use a different machine learning algorithm If any of these reduces the RMSE significantly, you have succeeded in improving your model!

```
#The predict() function in R is used to predict the values based on the input data.  
predictions <- predict(model, test.data)
```

```
# assign test.data to a new object, so that we can make changes  
result <- test.data
```

```
#add the new variable of predictions to the result object  
result <- cbind(result, predictions)
```

```
#add the results to a data frame containing test data and the prediction  
result <- cbind(g[row.names(result), ], predictions)
```

Making predictions

```
result %>%  
  ggplot() +
```

```
geom_point(aes(x = predictions, y = max_WL, color = Parasite_challenge, size = delta)) +
geom_abline() +
labs(x = "Predictions: Maximum weight loss", y = "Observed: Maximum weight loss",
     title = "Predicting tolerance, Weight loss in response to immune gene expression") +
theme_bw()
```

Visualizations

Warning: Removed 3 rows containing missing values (geom_point).



```
g$current_infection <- as.factor(g$current_infection)

#now select the genes and the actual infection of the mice in the new mutate column
#infection
g.2 <- g %>%
  dplyr::select(c(current_infection, delta, all_of(Genes)))

# to get reproducible results we use a seed
set.seed(42)

# We want the current infection to be predicted by the data in all of the other columns
# iter = how many random forests are needed, in theory 6 are enough
```



```
#now we can impute our data
g.imputed_parasite <- rfImpute(current_infection ~ ., data = g.2, iter = 6)
```

Predicting eimeria species according to gene expression

```
## ntree      OOB      1      2      3      4
##   300:  34.48% 40.91% 33.33%100.00% 19.57%
## ntree      OOB      1      2      3      4
##   300:  35.34% 45.45% 33.33%100.00% 19.57%
## ntree      OOB      1      2      3      4
##   300:  40.52% 50.00% 35.90%100.00% 28.26%
## ntree      OOB      1      2      3      4
##   300:  36.21% 40.91% 33.33%100.00% 23.91%
## ntree      OOB      1      2      3      4
##   300:  36.21% 36.36% 33.33%100.00% 26.09%
## ntree      OOB      1      2      3      4
##   300:  36.21% 40.91% 35.90%100.00% 21.74%
```

```
g.imputed_parasite <- g.imputed_parasite %>% dplyr::select(- current_infection)

g_minus <- g %>% dplyr::select(-c((all_of(Genes)), delta))

#full data set containing the imputed gene expression data
g.imputed_parasite <- cbind(g_minus, g.imputed_parasite)
```

```
g.imputed_parasite$current_infection <- as.factor(g.imputed_parasite$current_infection)

#select the relevant columns:
g.imputed_parasite <- g.imputed_parasite %>%
  dplyr::select(c(current_infection, delta, all_of(Genes)))

# split data into training and test
set.seed(123) # this will help us reproduce this random assignment
# in this way we can pick the random numbers
training.samples_parasite <- g.imputed_parasite$current_infection%>%
  createDataPartition(p = .7, # this is the partiicipation! In this case 0.7 = training data and 0.3 = test data
    list = FALSE) # we don't want to get a list in return
train.data_parasite <- g.imputed_parasite[training.samples, ] #include all the randomly selected rows
test.data_parasite <- g.imputed_parasite[-training.samples, ]
```

Now split the data again into training and testing

```
#train the model
model_Parasite <- randomForest(current_infection ~., data = train.data_parasite, proximity = TRUE,
  ntree = 1500) # number of trees

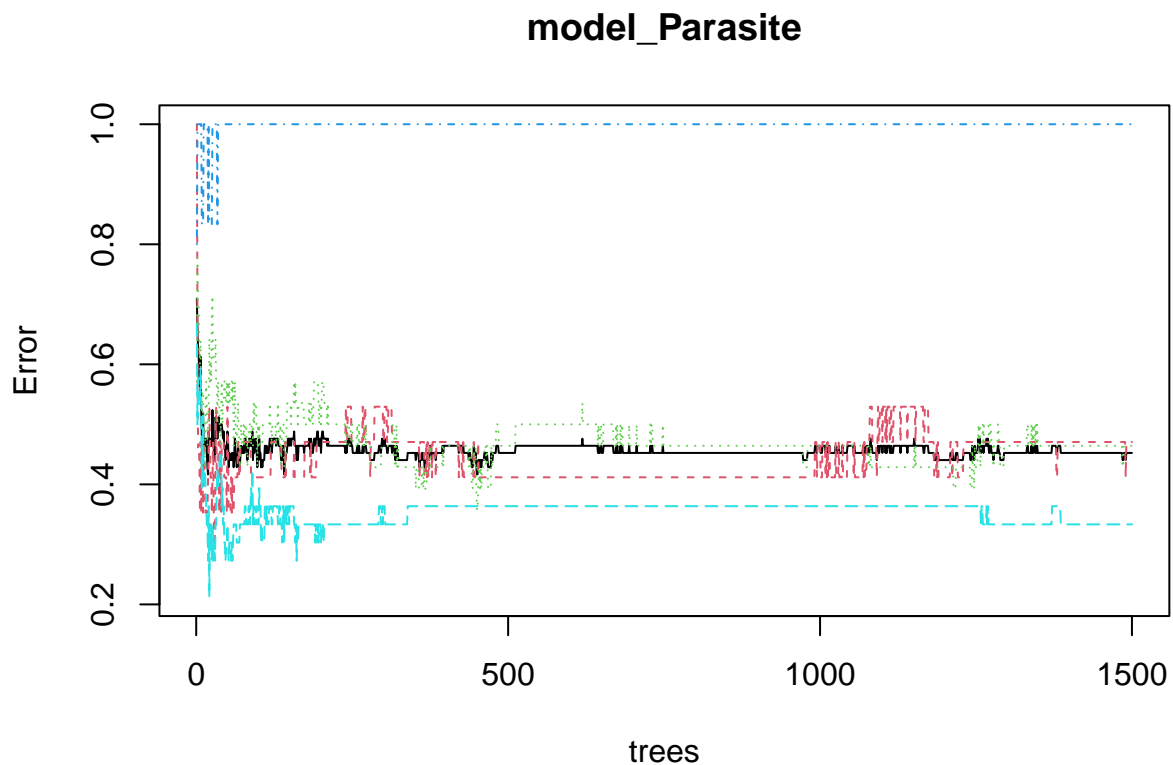
print(model_Parasite)
```

Building the model

```
##
## Call:
## randomForest(formula = current_infection ~ ., data = train.data_parasite,      proximity = TRUE, nt
##           Type of random forest: classification
##           Number of trees: 1500
## No. of variables tried at each split: 4
##
##           OOB estimate of  error rate: 45.24%
## Confusion matrix:
##           E_falciiformis E_ferrisi infected_eimeria uninfected
## E_falciiformis           9         6              0          2
## E_ferrisi                6        15              0          7
## infected_eimeria         0         1              0          5
## uninfected               1        10              0         22
##
##           class.error
## E_falciiformis  0.4705882
## E_ferrisi       0.4642857
## infected_eimeria 1.0000000
## uninfected      0.3333333
```

OOB = 46.43, this means that only 53 % of our predictions are accurate

```
plot(model_Parasite)
```



Test the model

```
#The predict() function in R is used to predict the values based on the input data.
predictions_parasite <- predict(model_Parasite, test.data_parasite)
# assign test.data to a new object, so that we can make changes
result_parasite <- test.data_parasite
#add the new variable of predictions to the result object
result_parasite <- cbind(result_parasite, predictions_parasite)
#add the results to a data frame containing test data and the prediction
result_parasite <- cbind(g[row.names(result_parasite), ], predictions_parasite)
```

Making predictions same but for parasite challenge instead of current infection

```
g$Parasite_challenge <- as.factor(g$Parasite_challenge)

#now select the genes and the actual infection of the mice in the new mutate column
#infection
g.2 <- g %>%
  dplyr::select(c(Parasite_challenge, delta, all_of(Genes)))

# to get reproducible results we use a seed
set.seed(42)
# We want the current infection to be predicted by the data in all of the other columns
# iter = how many random forests are needed, in theory 6 are enough

#now we can impute our data
g.imputed_parasite <- rfImpute(Parasite_challenge ~ ., data = g.2, iter = 6)
```

Predicting eimeria species according to gene expression

```
## ntree      OOB      1      2      3
##   300:  31.03% 54.55% 31.91% 19.15%
## ntree      OOB      1      2      3
##   300:  30.17% 50.00% 29.79% 21.28%
## ntree      OOB      1      2      3
##   300:  27.59% 45.45% 29.79% 17.02%
## ntree      OOB      1      2      3
##   300:  27.59% 50.00% 27.66% 17.02%
## ntree      OOB      1      2      3
##   300:  29.31% 45.45% 31.91% 19.15%
## ntree      OOB      1      2      3
##   300:  27.59% 50.00% 27.66% 17.02%
```

```
g.imputed_parasite <- g.imputed_parasite %>% dplyr::select(- Parasite_challenge)

g_minus <- g %>% dplyr::select(-c((all_of(Genes)), delta))
```

```

#full data set containing the imputed gene expression data
g.imputed_parasite <- cbind(g_minus, g.imputed_parasite)

g.imputed_parasite$Parasite_challenge <- as.factor(g.imputed_parasite$Parasite_challenge)

#select the relevant columns:
g.imputed_parasite <- g.imputed_parasite %>%
  dplyr::select(c(Parasite_challenge, Eim_MC, all_of(Genes)))

# to use in the next model
parasite_data <- g.imputed_parasite

g.imputed_parasite <- g.imputed_parasite %>%
  dplyr::select(-Eim_MC)

# split data into training and test
set.seed(123) # this will help us reproduce this random assignment
# in this way we can pick the random numbers
training.samples_parasite <- g.imputed_parasite$Parasite_challenge%>%
  createDataPartition(p = .7, # this is the partiicipation! In this case 0.7 = training data and 0.3 = te
    list = FALSE) # we don't want to get a list in return
train.data_parasite <- g.imputed_parasite[training.samples, ] #include all the randomly selected rows
test.data_parasite <- g.imputed_parasite[-training.samples, ]

```

Now split the data again into training and testing

```

#train the model
model_Parasite <- randomForest(Parasite_challenge ~ ., data = train.data_parasite, proximity = TRUE,
  ntree = 1500) # number of trees

print(model_Parasite)

```

Building the model

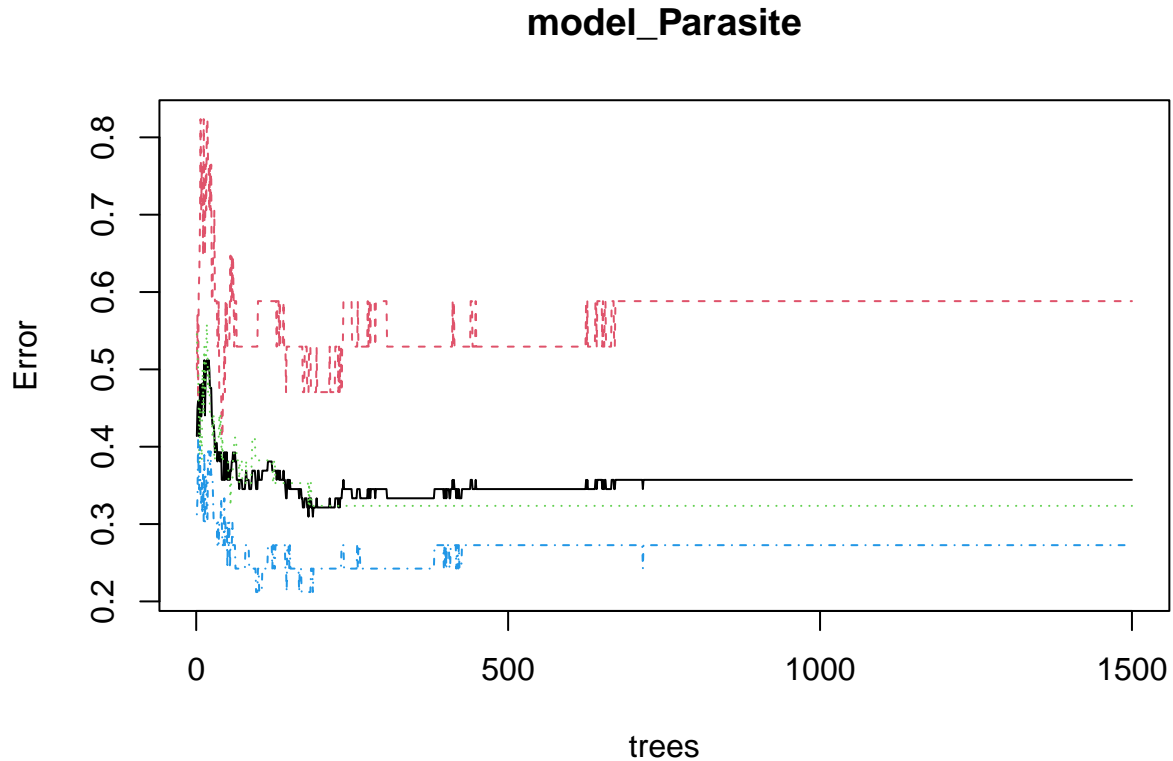
```

##
## Call:
## randomForest(formula = Parasite_challenge ~ ., data = train.data_parasite, proximity = TRUE, n
##           Type of random forest: classification
##           Number of trees: 1500
## No. of variables tried at each split: 4
##
##           OOB estimate of  error rate: 35.71%
## Confusion matrix:
##           E_falciformis E_ferrisi uninfected class.error
## E_falciformis          7          7          3  0.5882353
## E_ferrisi             4         23          7  0.3235294
## uninfected            1          8         24  0.2727273

```

OOB = 46.43, this means that only 53 % of our predictions are accurate

```
plot(model_Parasite)
```



Test the model

```
#The predict() function in R is used to predict the values based on the input data.
predictions_parasite <- predict(model_Parasite, test.data_parasite)
# assign test.data to a new object, so that we can make changes
result_parasite <- test.data_parasite
#add the new variable of predictions to the result object
result_parasite <- cbind(result_parasite, predictions_parasite)
#add the results to a data frame containing test data and the prediction
result_parasite <- cbind(g[row.names(result_parasite), ], predictions_parasite)
```

Making predictions

```
conf_matrix_parasite <- confusionMatrix(result_parasite$predictions_parasite, reference = result_parasite)
print(conf_matrix_parasite)
```

Visualizations


```
## Confusion Matrix and Statistics
##
##               Reference
## Prediction      E_falciformis E_ferrisi uninfected
##   E_falciformis           1         2         1
##   E_ferrisi              3         8         1
##   uninfected             1         3        12
##
## Overall Statistics
##
##               Accuracy : 0.6562
##               95% CI : (0.4681, 0.8143)
##   No Information Rate : 0.4375
##   P-Value [Acc > NIR] : 0.01044
##
##               Kappa : 0.4359
##
##   McNemar's Test P-Value : 0.75300
##
## Statistics by Class:
##
##               Class: E_falciformis Class: E_ferrisi Class: uninfected
## Sensitivity                0.20000         0.6154         0.8571
## Specificity                0.88889         0.7895         0.7778
## Pos Pred Value              0.25000         0.6667         0.7500
## Neg Pred Value              0.85714         0.7500         0.8750
## Prevalence                  0.15625         0.4062         0.4375
## Detection Rate              0.03125         0.2500         0.3750
## Detection Prevalence        0.12500         0.3750         0.5000
## Balanced Accuracy           0.54444         0.7024         0.8175
```

```
conf_matrix_parasite$table
```

```
##               Reference
## Prediction      E_falciformis E_ferrisi uninfected
##   E_falciformis           1         2         1
##   E_ferrisi              3         8         1
##   uninfected             1         3        12
```

```
plt <- as.data.frame(conf_matrix_parasite$table)
plt$Prediction <- factor(plt$Prediction, levels=rev(levels(plt$Prediction)))

ggplot(plt, aes(x = Prediction, y = Reference, fill= Freq)) +
  geom_tile() + geom_text(aes(label=Freq)) +
  scale_fill_gradient(low="white", high="forestgreen") +
  labs(x = "Predictions", y = "Reference")
```



```
train.data_parasite %>%
  group_by(Parasite_challenge) %>%
  summarize(length(Parasite_challenge))
```

```
## # A tibble: 3 x 2
##   Parasite_challenge 'length(Parasite_challenge)'
##   <fct>              <int>
## 1 E_falciformis      17
## 2 E_ferrisi         34
## 3 uninfected        33
```

Repeat the previous model, this time testing for

infected with Eimeria or not. #####

to use in the next model

```
parasite_data <- parasite_data %>%
  dplyr::select(-Parasite_challenge)
```

split data into training and test

```
set.seed(123) # this will help us reproduce this random assignment
```

in this way we can pick the random numbers

```
training.samples_melting <- parasite_data$Eim_MC%>%
```

```
  createDataPartition(p = .7, # this is the partiicipation! In this case 0.7 = training data and 0.3 = te
```

```

      list = FALSE) # we don't want to get a list in return
train.data_melting <- parasite_data[training.samples, ] #include all the randomly selected rows
test.data_melting <- parasite_data[-training.samples, ]

```

```

#train the model
model_melting <- randomForest(Eim_MC ~., data = train.data_melting, proximity = TRUE,
                             ntree = 1500) # number of trees

print(model_melting)

```

Building the model

```

##
## Call:
##  randomForest(formula = Eim_MC ~ ., data = train.data_melting,      proximity = TRUE, ntree = 1500)
##              Type of random forest: classification
##              Number of trees: 1500
## No. of variables tried at each split: 4
##
##      OOB estimate of  error rate: 30.95%
## Confusion matrix:
##      FALSE TRUE class.error
## FALSE    16   17   0.5151515
## TRUE      9   42   0.1764706

```

Test the model

```

#The predict() function in R is used to predict the values based on the input data.
predictions_melting <- predict(model_melting, test.data_melting)
# assign test.data to a new object, so that we can make changes
result_melting <- test.data_melting
#add the new variable of predictions to the result object
result_melting <- cbind(result_melting, predictions_melting)
#add the results to a data frame containing test data and the prediction
result_melting <- cbind(g[row.names(result_melting), ], predictions_melting)

```

```

conf_matrix_melting <- confusionMatrix(result_melting$predictions_melting, reference = result_melting$Eim_MC)
print(conf_matrix_melting)

```

Making predictions

```

## Confusion Matrix and Statistics
##

```



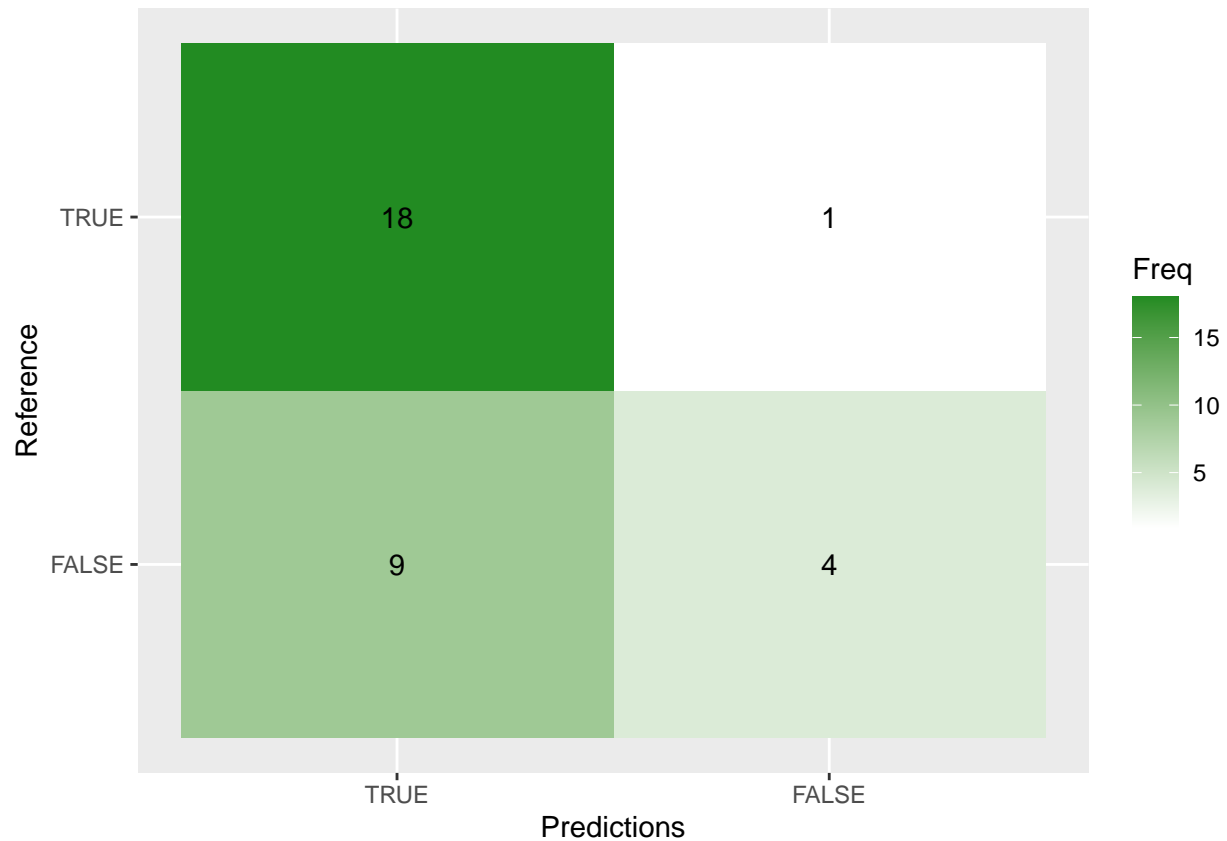
```
##           Reference
## Prediction FALSE TRUE
##      FALSE      4      1
##      TRUE       9     18
##
##           Accuracy : 0.6875
##           95% CI : (0.4999, 0.8388)
##      No Information Rate : 0.5938
##      P-Value [Acc > NIR] : 0.18482
##
##           Kappa : 0.2825
##
## Mcnemar's Test P-Value : 0.02686
##
##           Sensitivity : 0.3077
##           Specificity : 0.9474
##      Pos Pred Value : 0.8000
##      Neg Pred Value : 0.6667
##           Prevalence : 0.4062
##      Detection Rate : 0.1250
##      Detection Prevalence : 0.1562
##      Balanced Accuracy : 0.6275
##
##      'Positive' Class : FALSE
##
```

```
conf_matrix_melting$table
```

```
##           Reference
## Prediction FALSE TRUE
##      FALSE      4      1
##      TRUE       9     18
```

```
plt <- as.data.frame(conf_matrix_melting$table)
plt$Prediction <- factor(plt$Prediction, levels=rev(levels(plt$Prediction)))

ggplot(plt, aes(x = Prediction, y = Reference, fill= Freq)) +
  geom_tile() + geom_text(aes(label=Freq)) +
  scale_fill_gradient(low="white", high="forestgreen") +
  labs(x = "Predictions", y = "Reference")
```



Field data

```
Field <- read.csv("https://raw.githubusercontent.com/derele/Mouse_Eimeria_Field/master/data_products/SO")
```

Importing field data

```
Field %>% summarise(length(Mouse_ID))
```

Summary statistics for the field data

```
## length(Mouse_ID)
## 1 1921
```

We have 1921 mice in total.

```
EqPCR.cols <- c("delta_ct_cewe_MminusE", "MC.Eimeria", "Ct.Eimeria") #,"Ct.Mus""delta_ct_ilwe_MminusE"
EimGeno.cols <- c("n18S_Seq", "COI_Seq", "ORF470_Seq", "eimeriaSpecies")
```

```
Gene.Exp.cols <- c("IFNy", "CXCR3", "IL.6", #"GBP2", "IL.12", "IRG6",
                  "IL.10", "IL.13", "IL.10", "IL.13", "IL1RN",
                  "CXCR3", "CASP1", "CXCL9",
                  "IDO1", "IRGM1", "MPO", "MUC2", "MUC5AC", "MYD88",
                  "NCR1", "PRF1", "RETNLB", "SOCS1", "TICAM1", "TNF")
```

```
House.Keeping.cols <- c("GAPDH", "PPIB", "B.actin", "B-actin")
```

```
#which are the numbers of the columns of Field
names <- data.frame(colnames(Field))
```

```
f <- Field[ , c(76:78, 80:97)]
```

```
#how many nas in each column
sapply(f, function(x) sum(is.na(x)))
```

```
## CASP1 CXCL9 CXCR3 IDO1 IFNy IL.10 IL.13 IL1RN IRGM1 MPO MUC2
## 1713 1627 1695 1614 1592 1802 1613 1616 1596 1624 1599
## MUC5AC MYD88 NCR1 PPIB PRF1 RETNLB SOCS1 TICAM1 TNF IL.6
## 1615 1605 1711 1797 1720 1693 1596 1703 1625 1684
```

```
#remove rows with only nas
f <- f[rowSums(is.na(f)) != ncol(f), ]
```

```
Field <- Field %>%
  dplyr::select(-c(76:78, 80:97))
```

```
#merge the data frame to keep only the selected rows
f <- merge(Field, f, by = "row.names")
```

Imputing missing data

For the lab data I have used the function `rfimpute` from the package `random forest`. I can't use the same function for our lab data as the function requires the data set to contain predictor variable and response variables.

Therefore I will be using the package `MICE` (multivariate Imputation by chained Equations) which only requires a data frame of missing observations.

Description: *Multiple imputation using Fully Conditional Specification (FCS)*

implemented by the `MICE` algorithm as described in Van Buuren and Groothuis-Oudshoorn (2011) doi:10.18637/jss.v045.i03. Each variable has its own imputation model. Built-in imputation models are provided for continuous data (predictive mean matching, normal), binary data (logistic regression), unordered categorical data (polytomous logistic regression) and ordered categorical data (proportional odds). `MICE` can also impute continuous two-level data (normal model, pan, second-level variables). Passive imputation can be used to maintain consistency between variables. Various diagnostic plots are available to inspect the quality of the imputations.

<https://www.jstatsoft.org/article/view/v045i03>

tutorial: <https://www.youtube.com/watch?v=WPiYOS3qK70>

<https://datascienceplus.com/imputing-missing-data-with-r-mice-package/>

<https://datascienceplus.com/handling-missing-data-with-mice-package-a-simple-approach/>

Missing data can be classified into three categories:

1. **Missing completely at random (MCAR)** We can't probably predict that value from any other value in the data. MCAR implies the reason for the missingness of a field is completely random, and that we probably can't predict that value from any other value in the data.
2. **Missing at Random (MAR)** Missingness can be explained by other values in other columns, but not from that column.
3. **Missing NOT at random (MNAR)** The basic MICE assumption is that the data is missing at random, and that we can make a guess about its true value by looking at other data samples.

```
library(mice)

f <- f %>% dplyr::select(-"Row.names")

#turn the eimeria species into logical
f$eimeriaSpecies <- as.factor(f$eimeriaSpecies)

field_genes <- f %>%
  dplyr::select(Gene.Exp.cols)
```

Let's start by cleaning and checking the missing data points in our field data.


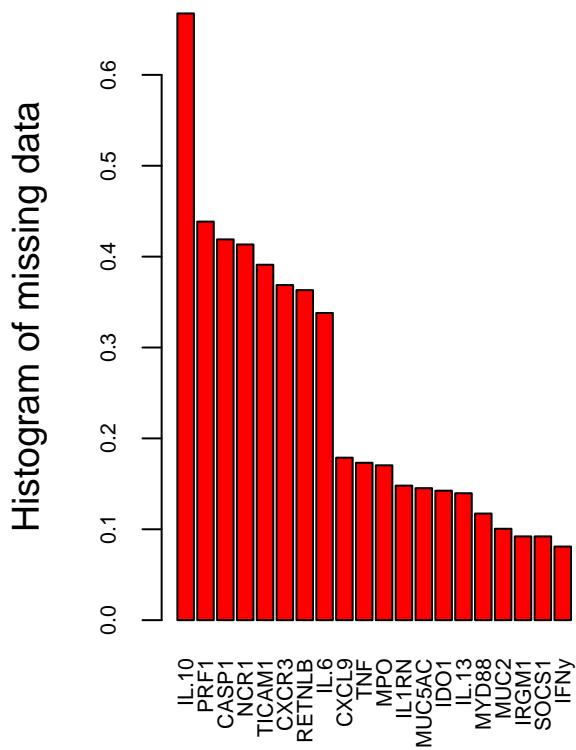
```
## Note: Using an external vector in selections is ambiguous.
## i Use 'all_of(Gene.Exp.cols)' instead of 'Gene.Exp.cols' to silence this message.
## i See <https://tidyselect.r-lib.org/reference/faq-external-vector.html>.
## This message is displayed once per session.
```

```
# check the data for missing values
sapply(field_genes, function(x) sum(is.na(x)))
```

```
##   IFNy   CXCR3   IL.6   IL.10   IL.13   IL1RN   CASP1   CXCL9   IDO1   IRGM1   MPO
##    29    132    121    239     50     53    150    64     51     33    61
##   MUC2 MUC5AC MYD88   NCR1   PRF1 RETNLB   SOCS1 TICAM1   TNF
##    36     52     42    148    157    130     33    140    62
```


```
field_genes %>%
  aggr(col = c('navyblue', 'red'), numbers = TRUE, sortVars = TRUE, labels=names(field_genes), cex.axis=
```

```
## Warning in plot.aggr(res, ...): not enough vertical space to display frequencies
## (too many combinations)
```



```
##
## Variables sorted by number of missings:
## Variable      Count
## IL.10 0.66759777
## PRF1 0.43854749
## CASP1 0.41899441
## NCR1 0.41340782
## TICAM1 0.39106145
## CXCR3 0.36871508
## RETNLB 0.36312849
## IL.6 0.33798883
## CXCL9 0.17877095
## TNF 0.17318436
## MPO 0.17039106
## IL1RN 0.14804469
## MUC5AC 0.14525140
## IDO1 0.14245810
## IL.13 0.13966480
## MYD88 0.11731844
## MUC2 0.10055866
## IRGM1 0.09217877
## SOCS1 0.09217877
## IFNy 0.08100559
```

```
marginplot(field_genes[c(1,2)])
```



```
# The frequency distribution of the missing cases per variable can be obtained as:
init <- mice(field_genes, maxit = 0)

# table of amount of variables with the amount of missing values
table(init$nmis)
```

Now let's continue by using the package MICE to impute the data

```
##
## 29 33 36 42 50 51 52 53 61 62 64 121 130 132 140 148 150 157 239
## 1 2 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```

```
# which method is used for imputation? In this case the package mice
# uses the default method for continuous variable,
# which is pmm, or predictive mean matching

meth <- init$method

# now impute the immune gene expression for the field and save it as the object:
```

```
# igf
# m=5 refers to the number of imputed datasets. Five is the default value.
igf <- mice(field_genes, method = meth, m = 5, seed = 500)
```

```
##
## iter imp variable
## 1 1 IFNy CXCR3 IL.6 IL.10 IL.13 IL1RN CASP1 CXCL9 ID01 IRGM1 MPO MUC2 MUC5AC MYD88
## 1 2 IFNy CXCR3 IL.6 IL.10 IL.13 IL1RN CASP1 CXCL9 ID01 IRGM1 MPO MUC2 MUC5AC MYD88
## 1 3 IFNy CXCR3 IL.6 IL.10 IL.13 IL1RN CASP1 CXCL9 ID01 IRGM1 MPO MUC2 MUC5AC MYD88
## 1 4 IFNy CXCR3 IL.6 IL.10 IL.13 IL1RN CASP1 CXCL9 ID01 IRGM1 MPO MUC2 MUC5AC MYD88
## 1 5 IFNy CXCR3 IL.6 IL.10 IL.13 IL1RN CASP1 CXCL9 ID01 IRGM1 MPO MUC2 MUC5AC MYD88
## 2 1 IFNy CXCR3 IL.6 IL.10 IL.13 IL1RN CASP1 CXCL9 ID01 IRGM1 MPO MUC2 MUC5AC MYD88
## 2 2 IFNy CXCR3 IL.6 IL.10 IL.13 IL1RN CASP1 CXCL9 ID01 IRGM1 MPO MUC2 MUC5AC MYD88
## 2 3 IFNy CXCR3 IL.6 IL.10 IL.13 IL1RN CASP1 CXCL9 ID01 IRGM1 MPO MUC2 MUC5AC MYD88
## 2 4 IFNy CXCR3 IL.6 IL.10 IL.13 IL1RN CASP1 CXCL9 ID01 IRGM1 MPO MUC2 MUC5AC MYD88
## 2 5 IFNy CXCR3 IL.6 IL.10 IL.13 IL1RN CASP1 CXCL9 ID01 IRGM1 MPO MUC2 MUC5AC MYD88
## 3 1 IFNy CXCR3 IL.6 IL.10 IL.13 IL1RN CASP1 CXCL9 ID01 IRGM1 MPO MUC2 MUC5AC MYD88
## 3 2 IFNy CXCR3 IL.6 IL.10 IL.13 IL1RN CASP1 CXCL9 ID01 IRGM1 MPO MUC2 MUC5AC MYD88
## 3 3 IFNy CXCR3 IL.6 IL.10 IL.13 IL1RN CASP1 CXCL9 ID01 IRGM1 MPO MUC2 MUC5AC MYD88
## 3 4 IFNy CXCR3 IL.6 IL.10 IL.13 IL1RN CASP1 CXCL9 ID01 IRGM1 MPO MUC2 MUC5AC MYD88
## 3 5 IFNy CXCR3 IL.6 IL.10 IL.13 IL1RN CASP1 CXCL9 ID01 IRGM1 MPO MUC2 MUC5AC MYD88
## 4 1 IFNy CXCR3 IL.6 IL.10 IL.13 IL1RN CASP1 CXCL9 ID01 IRGM1 MPO MUC2 MUC5AC MYD88
## 4 2 IFNy CXCR3 IL.6 IL.10 IL.13 IL1RN CASP1 CXCL9 ID01 IRGM1 MPO MUC2 MUC5AC MYD88
## 4 3 IFNy CXCR3 IL.6 IL.10 IL.13 IL1RN CASP1 CXCL9 ID01 IRGM1 MPO MUC2 MUC5AC MYD88
## 4 4 IFNy CXCR3 IL.6 IL.10 IL.13 IL1RN CASP1 CXCL9 ID01 IRGM1 MPO MUC2 MUC5AC MYD88
## 4 5 IFNy CXCR3 IL.6 IL.10 IL.13 IL1RN CASP1 CXCL9 ID01 IRGM1 MPO MUC2 MUC5AC MYD88
## 5 1 IFNy CXCR3 IL.6 IL.10 IL.13 IL1RN CASP1 CXCL9 ID01 IRGM1 MPO MUC2 MUC5AC MYD88
## 5 2 IFNy CXCR3 IL.6 IL.10 IL.13 IL1RN CASP1 CXCL9 ID01 IRGM1 MPO MUC2 MUC5AC MYD88
## 5 3 IFNy CXCR3 IL.6 IL.10 IL.13 IL1RN CASP1 CXCL9 ID01 IRGM1 MPO MUC2 MUC5AC MYD88
## 5 4 IFNy CXCR3 IL.6 IL.10 IL.13 IL1RN CASP1 CXCL9 ID01 IRGM1 MPO MUC2 MUC5AC MYD88
## 5 5 IFNy CXCR3 IL.6 IL.10 IL.13 IL1RN CASP1 CXCL9 ID01 IRGM1 MPO MUC2 MUC5AC MYD88
```

```
summary(igf)
```

```
## Class: mids
## Number of multiple imputations: 5
## Imputation methods:
## IFNy CXCR3 IL.6 IL.10 IL.13 IL1RN CASP1 CXCL9 ID01 IRGM1 MPO
## "pmm" "pmm" "pmm" "pmm" "pmm" "pmm" "pmm" "pmm" "pmm" "pmm" "pmm"
## MUC2 MUC5AC MYD88 NCR1 PRF1 RETNLB SOCS1 TICAM1 TNF
## "pmm" "pmm" "pmm" "pmm" "pmm" "pmm" "pmm" "pmm" "pmm"
## PredictorMatrix:
## IFNy CXCR3 IL.6 IL.10 IL.13 IL1RN CASP1 CXCL9 ID01 IRGM1 MPO MUC2 MUC5AC
## IFNy 0 1 1 1 1 1 1 1 1 1 1 1 1
## CXCR3 1 0 1 1 1 1 1 1 1 1 1 1 1
## IL.6 1 1 0 1 1 1 1 1 1 1 1 1 1
## IL.10 1 1 1 0 1 1 1 1 1 1 1 1 1
## IL.13 1 1 1 1 0 1 1 1 1 1 1 1 1
## IL1RN 1 1 1 1 1 0 1 1 1 1 1 1 1
## MYD88 NCR1 PRF1 RETNLB SOCS1 TICAM1 TNF
## IFNy 1 1 1 1 1 1 1
## CXCR3 1 1 1 1 1 1 1
```

```
## IL.6      1    1    1    1    1    1    1
## IL.10     1    1    1    1    1    1    1
## IL.13     1    1    1    1    1    1    1
## IL1RN     1    1    1    1    1    1    1
```

```
# to check each column with imputed data
```

```
## igf$imp$IFNy
```


```
#Now we can get back the completed dataset using the complete()
```

```
completeField <- complete(igf, 1)
```

Predictive mean matching with $d = 5$ is the default in `mice()` for continuous data. The method is robust against misspecification of the imputation model, yet performs as well as theoretically superior methods. In the context of missing covariate data, Marshall, Altman, and Holder (2010) concluded that predictive mean matching “produced the least biased estimates and better model performance measures.” Another simulation study that addressed skewed data concluded that predictive mean matching “may be the preferred approach provided that less than 50% of the cases have missing data and the missing data are not MNAR” (Marshall et al. 2010). Kleinke (2017) found that the method works well across a wide variety of scenarios, but warned the default cannot address severe skewness or small samples.

Let’s compare the distributions of original and imputed data using a some useful plots. First of all we can use a scatterplot and plot Ozone against all the other variables Let’s first plot the variables for which we have few missing values


```
xyplot(igf, IFNy ~ SOCS1 + IRGM1 + MUC2, pch=18, cex=1)
```



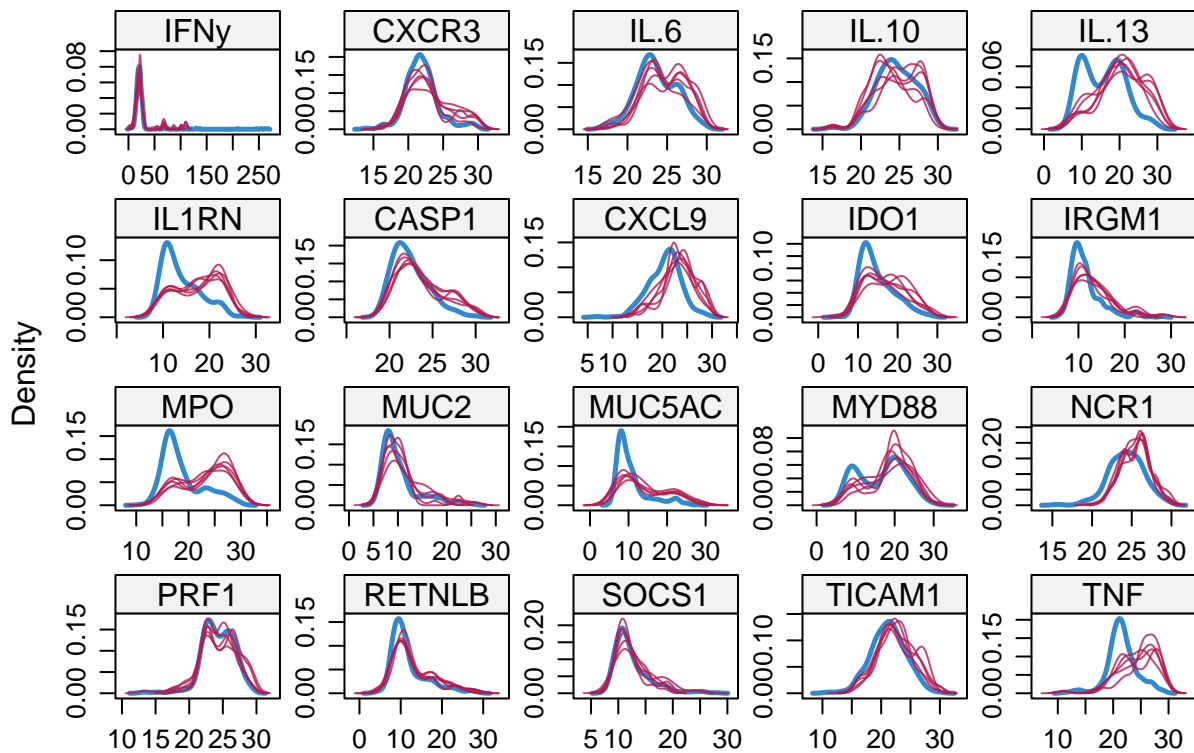
What we would like to see is that the shape of the magenta points (imputed) matches the shape of the blue ones (observed). The matching shape tells us that the imputed values are indeed “plausible values”.

Now let’s plot the variables with many missing data points.

```
xyplot(igf, IFNy ~ IL.10 + PRF1 + CASP1, pch=18, cex=1)
```



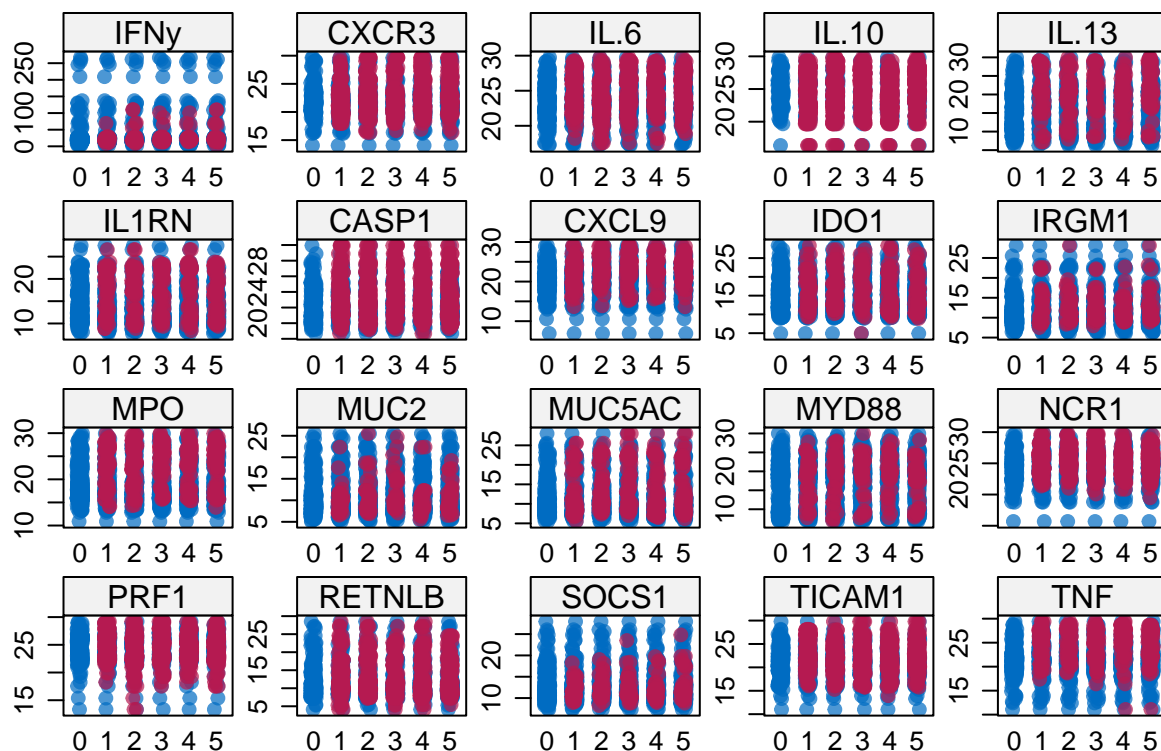
```
densityplot(igf)
```



The density of the imputed data for each imputed dataset is showed in magenta while the density of the observed data is showed in blue. Again, under our previous assumptions we expect the distributions to be similar.

Another useful visual take on the distributions can be obtained using the `stripplot()` function that shows the distributions of the variables as individual points

```
stripplot(igf, pch = 20, cex = 1.2)
```



Applying the model for predicting weight loss to our imputed data set Start by making the predictions for the field data.

```
# Start by selecting the columns that appear in both the training data set and the
# field data set
completeField <- completeField %>%
  dplyr::select(intersect(colnames(completeField), colnames(train.data)))

set.seed(540)

#The predict() function in R is used to predict the values based on the input data.
predictions_field <- predict(model, completeField)

# assign test.data to a new object, so that we can make changes
result_field <- completeField


#add the new variable of predictions to the result object
result_field <- cbind(result_field, predictions_field)

#add the results to a data frame containing test data and the prediction
f <- f %>%
  dplyr::select(-intersect(colnames(result_field), colnames(f)))

result_field <- cbind(f, result_field)
```

It is time to apply the package of Alice Balard et al. on our predictions! Let's see if we indeed have differences across the hybrid index with our predicted weight loss.

```
result_field %>% ggplot(aes(x = predictions_field)) +  
  geom_histogram(binwidth = 1.5)
```




Check the distribution

```
result_field %>%  
  ggplot(aes(x = HI , y = predictions_field , color = Sex)) +  
  geom_smooth() +  
  geom_point()
```

```
## 'geom_smooth()' using method = 'loess' and formula 'y ~ x'
```

```
## Warning: Removed 1 rows containing non-finite values (stat_smooth).
```

```
## Warning: Removed 1 rows containing missing values (geom_point).
```




```
result_field %>%
  ggplot(aes(x = Body_Length , y = predictions_field , color = Sex)) +
  geom_smooth() +
  geom_point()
```

```
## 'geom_smooth()' using method = 'loess' and formula 'y ~ x'
```

```
## Warning: Removed 1 rows containing non-finite values (stat_smooth).
```

```
## Removed 1 rows containing missing values (geom_point).
```



Nice to see that they are normally distributed.

Fitting distributions??

Ratios / Percentages are not normally distributed. Weibull is a good distributions.

Alice used weibull for the qpcr data. (paper)

```
library(fitdistrplus)
```

```
## Loading required package: MASS
```

```
##
```

```
## Attaching package: 'MASS'
```

```
## The following object is masked from 'package:dplyr':
```

```
##
```

```
##      select
```

```
## Loading required package: survival
```

```
##
```

```
## Attaching package: 'survival'
```

```
## The following object is masked from 'package:caret':
```

```
##
```

```
##      cluster
```


```
library(logspline)

result_field <- result_field %>%
  dplyr::mutate(WL = predictions_field * (-1))

x <- result_field$WL

descdist(data = x, discrete = FALSE)
```


Cullen and Frey graph



```
## summary statistics
## -----
## min: 5.167267 max: 15.86692
## median: 9.722013
## mean: 10.15507
## estimated sd: 2.332844
## estimated skewness: 0.2244729
## estimated kurtosis: 2.041805
```

```
descdist(data = x, discrete = FALSE, #data is continuous
          boot = 1000)
```

Cullen and Frey graph



```
## summary statistics
## -----
## min: 5.167267 max: 15.86692
## median: 9.722013
## mean: 10.15507
## estimated sd: 2.332844
## estimated skewness: 0.2244729
## estimated kurtosis: 2.041805
```

Test for binomial distribution

```
set.seed(10)
n = 25
size = 27
prob = .4
data = rbinom(x, size = size, prob = prob)
fit = fitdist(data = data, dist="binom",
              fix.arg=list(size = size),
              start=list(prob = 0.1))

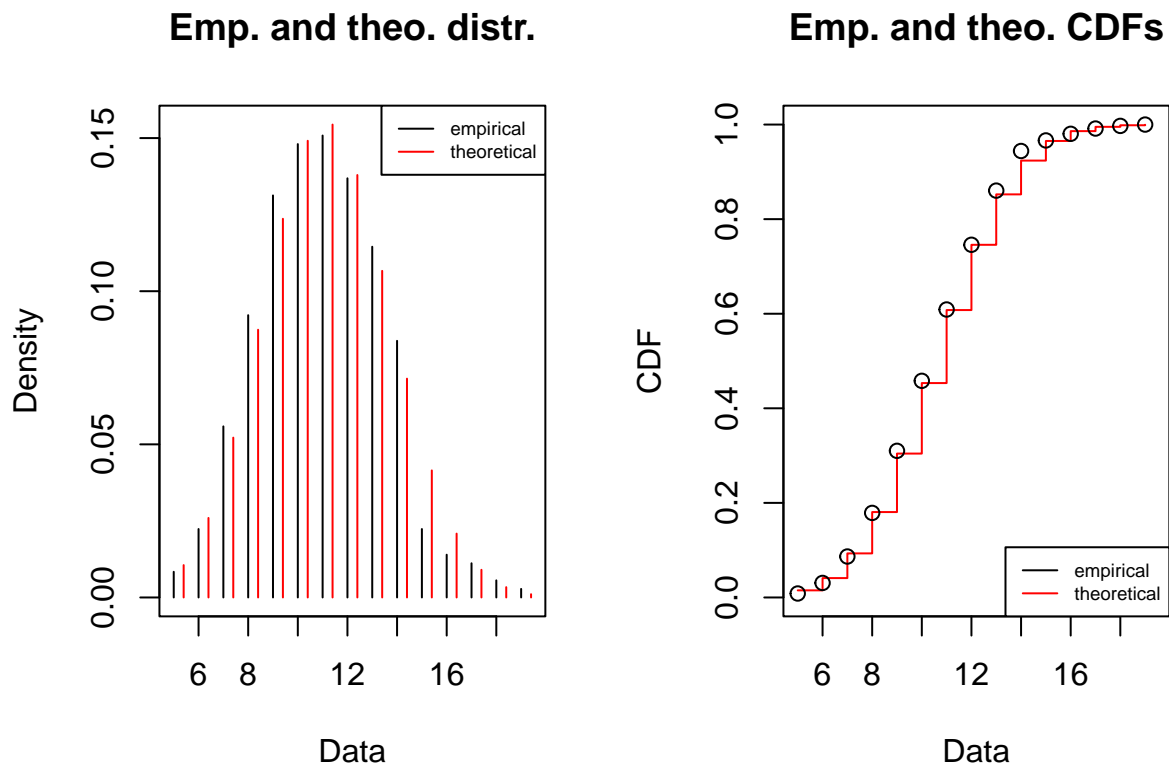
summary(fit)
```

```
## Fitting of the distribution ' binom ' by maximum likelihood
## Parameters :
##      estimate Std. Error
```



```
## prob 0.4012002 0.004985353
## Fixed parameters:
##      value
## size    27
## Loglikelihood: -832.1786   AIC: 1666.357   BIC: 1670.238
```

```
plot(fit)
```



```
normal_ <- fitdist(x, "norm")
weibull_ <- fitdist(x, "weibull")
gamma_ <- fitdist(x, "gamma")
```

```
library(fitdistrplus) # evaluate distribution

# Define function to be used to test, get the log lik and aic
tryDistrib <- function(x, distrib){
  # deals with fitdistr error:
  fit <- tryCatch(MASS::fitdistr(x, distrib), error=function(err) "fit failed")
  return(list(fit = fit,
              loglik = tryCatch(fit$loglik, error=function(err) "no loglik computed"),
              AIC = tryCatch(fit$aic, error=function(err) "no aic computed")))
}
```

```

findGoodDist <- function(x, distribs, distribs2){
  l =lapply(distribs, function(i) tryDistrib(x, i))
  names(l) <- distribs
  print(l)
  listDistr <- lapply(distribs2, function(i){
    if (i %in% "t"){
      fitdistrplus::fitdist(x, i, start = list(df =2))
    } else {
      fitdistrplus::fitdist(x,i)
    }
  })
  )
  par(mfrow=c(2,2))
  denscomp(listDistr, legendtext=distribs2)
  cdfcomp(listDistr, legendtext=distribs2)
  qqcomp(listDistr, legendtext=distribs2)
  ppcomp(listDistr, legendtext=distribs2)
  par(mfrow=c(1,1))
}

```

```
tryDistrib(x, "normal")
```

```

## $fit
##      mean      sd
## 10.15506798 2.32958395
## ( 0.12312234) ( 0.08706064)
##
## $loglik
## [1] -810.7369
##
## $AIC
## NULL

```

```
tryDistrib(x, "binomial")
```

```

## $fit
## [1] "fit failed"
##
## $loglik
## [1] "no loglik computed"
##
## $AIC
## [1] "no aic computed"

```

```
tryDistrib(x, "student")
```

```

## $fit
## [1] "fit failed"
##
## $loglik
## [1] "no loglik computed"
##
## $AIC
## [1] "no aic computed"

```

```
tryDistrib(x, "weibull")
```

```
## Warning in densfun(x, parm[1], parm[2], ...): NaNs produced
```

```
## Warning in densfun(x, parm[1], parm[2], ...): NaNs produced
```

```
## $fit
##      shape      scale
##  4.8050354 11.0891250
## ( 0.1964323) ( 0.1290913)
##
## $loglik
## [1] -813.8475
##
## $AIC
## NULL
```


```
tryDistrib(x, "weibullshifted")
```

```
## $fit
## [1] "fit failed"
##
## $loglik
## [1] "no loglik computed"
##
## $AIC
## [1] "no aic computed"
```


```
findGoodDist(x, "normal", "weibull")
```

```
## $normal
## $normal$fit
##      mean      sd
## 10.15506798 2.32958395
## ( 0.12312234) ( 0.08706064)
##
## $normal$loglik
## [1] -810.7369
##
## $normal$AIC
## NULL
```


Histogram and theoretical densities




Empirical and theoretical CDFs



Q-Q plot




P-P plot




```
plot(normal_)
```


Empirical and theoretical dens.




Q-Q plot



Empirical and theoretical CDFs



P-P plot




```
summary(normal_)
```


```
## Fitting of the distribution ' norm ' by maximum likelihood
## Parameters :
##      estimate Std. Error
## mean 10.155068 0.12312234
## sd   2.329584 0.08706057
## Loglikelihood: -810.7369   AIC: 1625.474   BIC: 1633.235
## Correlation matrix:
##      mean sd
## mean  1  0
## sd    0  1
```

```
plot(gamma_)
```


Empirical and theoretical dens.




Q-Q plot



Empirical and theoretical CDFs







P-P plot



```
summary(gamma_)
```

```
## Fitting of the distribution ' gamma ' by maximum likelihood
## Parameters :
##      estimate Std. Error
## shape 18.777962  1.3912338
## rate   1.849093  0.1388406
## Loglikelihood: -806.4038   AIC:  1616.808   BIC:  1624.569
## Correlation matrix:
##      shape      rate
## shape 1.0000000  0.9867194
## rate   0.9867194  1.0000000
```

```
plot(weibull_)
```

Empirical and theoretical dens.**Q-Q plot****Empirical and theoretical CDFs****P-P plot**

```
summary(weibull_)
```

```
## Fitting of the distribution ' weibull ' by maximum likelihood
## Parameters :
##      estimate Std. Error
## shape  4.804773  0.1964228
## scale 11.089299  0.1291004
## Loglikelihood: -813.8475   AIC:  1631.695   BIC:  1639.456
## Correlation matrix:
##      shape      scale
## shape 1.0000000 0.3273439
## scale 0.3273439 1.0000000
```

Is alpha significant for each hypothesis?

H0: the expected load for the subspecies and between 2 groups is the same H1: the mean load across 2 groups is the same, but can differ across subspecies H2: the mean load across subspecies is the same, but can differ between the 2 groups H3: the mean load can differ both across subspecies and between 2 groups

```
result_field$Sex <- as.factor(result_field$Sex)

result_field <- result_field %>%
  drop_na(HI)

parasiteLoad::getParamBounds("weibull", data = result_field, response = "WL")
```

```
##      L1start      L1LB      L1UB      L2start      L2LB      L2UB
## 10.153594373 0.000000001 15.866922159 10.153594373 0.000000001 15.866922159
##   alphaStart    alphaLB    alphaUB myshapeStart    myshapeLB    myshapeUB
## 0.000000000 -5.000000000 5.000000000 1.000000000 0.000000001 5.000000000
```

```
speparam <- c(L1start = 10,
              L1LB = 1e-9,
              L1UB = 20,
              L2start = 10,
              L2LB = 1e-9,
              L2UB = 20,
              alphaStart = 0, alphaLB = -5, alphaUB = 5,
              myshapeStart = 1, myshapeLB = 1e-9, myshapeUB = 5)
```

```
##All
```

```
parasiteLoad::analyse(data = result_field,
                      response = "WL",
                      model = "weibull",
                      group = "Sex")
```

```
## [1] "Analysing data for response: WL"
## [1] "Fit for the response: WL"
## [1] "Fitting for all"
## [1] "Fitting model basic without alpha"
## [1] "Did converge"
## [1] "Fitting model basic with alpha"
## [1] "Did converge"
## [1] "Fitting model advanced without alpha"
## [1] "Did converge"
## [1] "Fitting model advanced with alpha"
## [1] "Did converge"
## [1] "Fitting for groupA : F"
## [1] "Fitting model basic without alpha"
## [1] "Did converge"
## [1] "Fitting model basic with alpha"
## [1] "Did converge"
## [1] "Fitting model advanced without alpha"
## [1] "Did converge"
## [1] "Fitting model advanced with alpha"
## [1] "Did converge"
## [1] "Fitting for groupB : M"
## [1] "Fitting model basic without alpha"
## [1] "Did converge"
## [1] "Fitting model basic with alpha"
## [1] "Did converge"
## [1] "Fitting model advanced without alpha"
## [1] "Did converge"
## [1] "Fitting model advanced with alpha"
## [1] "Did converge"
## [1] "Testing H0 no alpha vs alpha"
##      dLL dDF      pvalue
## 1 1.26    1 0.1121067
## [1] "Testing H1 no alpha vs alpha"
##      dLL dDF      pvalue
```



```

## 1 1.21    1 0.1203272
## [1] "Testing H2 groupA no alpha vs alpha"
##      dLL dDF      pvalue
## 1 0.06    1 0.7284464
## [1] "Testing H2 groupB no alpha vs alpha"
##      dLL dDF      pvalue
## 1 1.98    1 0.04635762
## [1] "Testing H3 groupA no alpha vs alpha"
##      dLL dDF      pvalue
## 1 0.33    1 0.4160676
## [1] "Testing H3 groupB no alpha vs alpha"
##      dLL dDF      pvalue
## 1 2.61    1 0.02237828
## [1] "Testing H1 vs H0"
##      dLL dDF      pvalue
## 1 0.01    1 0.8913727
## [1] "Testing H2 vs H0"
##      dLL dDF      pvalue
## 1 0.86    3 0.6343848
## [1] "Testing H3 vs H1"
##      dLL dDF      pvalue
## 1 3.56    4 0.1294703
## [1] "Testing H3 vs H2"
##      dLL dDF      pvalue
## 1 2.72    2 0.06616066

## $H0
##
## Call:
## bbmle::mle2(minuslogl = response ~ dweibull(shape = myshape,
##      scale = MeanLoad(L1, L1, alpha, HI)/gamma(1 + (1/myshape))),
##      start = start, method = config$method, optimizer = config$optimizer,
##      data = data, lower = c(L1 = paramBounds[["L1LB"]], alpha = paramBounds[["alphaLB"]],
##      myshape = paramBounds[["myshapeLB"]]), upper = c(L1 = paramBounds[["L1UB"]],
##      alpha = paramBounds[["alphaUB"]], myshape = paramBounds[["myshapeUB"]]),
##      control = config$control)
##
## Coefficients:
##           L1      alpha    myshape
## 9.8668586 -0.1284202  4.8126675
##
## Log-likelihood: -810.77
## Best method: bobyqa
##
## $H1
##
## Call:
## bbmle::mle2(minuslogl = response ~ dweibull(shape = myshape,
##      scale = MeanLoad(L1, L2, alpha, HI)/gamma(1 + (1/myshape))),
##      start = start, method = config$method, optimizer = config$optimizer,
##      data = data, lower = c(L1 = paramBounds[["L1LB"]], L2 = paramBounds[["L2LB"]],
##      alpha = paramBounds[["alphaLB"]], myshape = paramBounds[["myshapeLB"]]),
##      upper = c(L1 = paramBounds[["L1UB"]], L2 = paramBounds[["L2UB"]],
##      alpha = paramBounds[["alphaUB"]], myshape = paramBounds[["myshapeUB"]]),

```

```

##      control = config$control)
##
## Coefficients:
##      L1      L2      alpha      myshape
##  9.8461464  9.8879260 -0.1266397  4.8120492
##
## Log-likelihood: -810.76
## Best method: bobyqa
##
## $H2
## $H2$groupA
##
## Call:
## bbmle::mle2(minuslogl = response ~ dweibull(shape = myshape,
##      scale = MeanLoad(L1, L1, alpha, HI)/gamma(1 + (1/myshape))),
##      start = start, method = config$method, optimizer = config$optimizer,
##      data = data, lower = c(L1 = paramBounds[["L1LB"]], alpha = paramBounds[["alphaLB"]],
##      myshape = paramBounds[["myshapeLB"]]), upper = c(L1 = paramBounds[["L1UB"]],
##      alpha = paramBounds[["alphaUB"]], myshape = paramBounds[["myshapeUB"]]),
##      control = config$control)
##
## Coefficients:
##      L1      alpha      myshape
## 10.02463142 -0.03702305  4.78582955
##
## Log-likelihood: -403.82
## Best method: bobyqa
##
## $H2$groupB
##
## Call:
## bbmle::mle2(minuslogl = response ~ dweibull(shape = myshape,
##      scale = MeanLoad(L1, L1, alpha, HI)/gamma(1 + (1/myshape))),
##      start = start, method = config$method, optimizer = config$optimizer,
##      data = data, lower = c(L1 = paramBounds[["L1LB"]], alpha = paramBounds[["alphaLB"]],
##      myshape = paramBounds[["myshapeLB"]]), upper = c(L1 = paramBounds[["L1UB"]],
##      alpha = paramBounds[["alphaUB"]], myshape = paramBounds[["myshapeUB"]]),
##      control = config$control)
##
## Coefficients:
##      L1      alpha      myshape
##  9.6609835 -0.2479875  4.8717470
##
## Log-likelihood: -406.09
## Best method: bobyqa
##
## $H3
## $H3$groupA
##
## Call:
## bbmle::mle2(minuslogl = response ~ dweibull(shape = myshape,
##      scale = MeanLoad(L1, L2, alpha, HI)/gamma(1 + (1/myshape))),
##      start = start, method = config$method, optimizer = config$optimizer,

```

```

##      data = data, lower = c(L1 = paramBounds[["L1LB"]], L2 = paramBounds[["L2LB"]],
##      alpha = paramBounds[["alphaLB"]], myshape = paramBounds[["myshapeLB"]]),
##      upper = c(L1 = paramBounds[["L1UB"]], L2 = paramBounds[["L2UB"]],
##      alpha = paramBounds[["alphaUB"]], myshape = paramBounds[["myshapeUB"]]),
##      control = config$control)
##
## Coefficients:
##          L1          L2          alpha      myshape
## 10.23938519  9.64263827 -0.09630871  4.81591546
##
## Log-likelihood: -403
## Best method: bobyqa
##
## $H3$groupB
##
## Call:
## bbmle::mle2(minuslogl = response ~ dweibull(shape = myshape,
##      scale = MeanLoad(L1, L2, alpha, HI)/gamma(1 + (1/myshape))),
##      start = start, method = config$method, optimizer = config$optimizer,
##      data = data, lower = c(L1 = paramBounds[["L1LB"]], L2 = paramBounds[["L2LB"]],
##      alpha = paramBounds[["alphaLB"]], myshape = paramBounds[["myshapeLB"]]),
##      upper = c(L1 = paramBounds[["L1UB"]], L2 = paramBounds[["L2UB"]],
##      alpha = paramBounds[["alphaUB"]], myshape = paramBounds[["myshapeUB"]]),
##      control = config$control)
##
## Coefficients:
##          L1          L2          alpha      myshape
##  9.1198785  9.9236196 -0.2813268  4.9194178
##
## Log-likelihood: -404.2
## Best method: bobyqa

```