

## 10. Applying random forest on field data - gene

Fay

2022-11-04

### Aim:

- Applying the models established in the script: 9
- How are hybrid mice different to the parental species?

### Load necessary libraries:

```
#install.packages("optimx", version = "2021-10.12") # this package is required for  
#the parasite load package to work  
library(tidyverse)  
library(tidyr)  
library(dplyr)  
library(cowplot)  
library(randomForest)  
library(ggplot2)  
library(VIM) # visualizing missing data  
library(mice) # imputing missing data without predictors  
library(ggpubr)  
library(optimx)  
library(rfUtilities) # Implements a permutation test cross-validation for  
# Random Forests models  
library(mice) #imputations  
library(fitdistrplus) #testing distributions  
library(logspline)  
library(caret)
```

### Field data

#### Import field data

```
hm <- read.csv("output_data/2.imputed_MICE_data_set.csv")
```

#### Clean data

```
Field <- hm %>%  
  filter(origin == "Field") %>%  
  drop_na(HI)
```

We have 1921 mice in total.

```
EqPCR.cols      <- c("delta_ct_cewe_MminusE", "MC.Eimeria", "Ct.Eimeria") #, "Ct.Mus" "delta_ct_ilwe_Mmin
Genes_wild      <- c("IFNy", "CXCR3", "IL.6", "IL.13", "IL.10",
                    "IL1RN", "CASP1", "CXCL9", "IDO1", "IRGM1", "MPO",
                    "MUC2", "MUC5AC", "MYD88", "NCR1", "PRF1", "RETNLB", "SOCS1",
                    "TICAM1", "TNF") #, "IL.12", "IRG6")
```

Prepare vectors for selecting

Actual Cleaning

```
#select the imputed gene columns
gene <- Field %>%
  dplyr::select(c(Mouse_ID, "IFNy", "CXCR3", "IL.6", "IL.13", "IL.10",
                  "IL1RN", "CASP1", "CXCL9", "IDO1", "IRGM1", "MPO",
                  "MUC2", "MUC5AC", "MYD88", "NCR1", "PRF1", "RETNLB", "SOCS1",
                  "TICAM1", "TNF"))

genes <- gene %>%
  dplyr::select(-Mouse_ID)

#remove rows with only nas
genes <- genes[,colSums(is.na(genes))<nrow(genes)]

#remove columns with only nas
genes <- genes[rowSums(is.na(genes)) != ncol(genes), ]

# select the same rows from the gene data
gene <- gene[row.names(genes),]

# select the same rows from the field data
Field <- Field[row.names(genes),]
```

## Predicting weight loss in our imputed field data

Start by making the predictions for the field data.

```
# load predicting weight loss model
weight_loss_predict <- readRDS("r_scripts/models/predict_WL.rds")

set.seed(540)

#The predict() function in R is used to predict the values based on the input data.
predictions_field <- predict(weight_loss_predict, genes)

#make the vector positive so that the distributions further down work
predictions_field <- predictions_field * (-1)

# assign test.data to a new object, so that we can make changes
result_field <- genes
```

```
#add the new variable of predictions to the result object
result_field <- cbind(result_field, predictions_field)

# add it to the field data
Field <- cbind(Field, predictions_field)
```

## It is time to apply the package of Alice Balard et al. on our predictions!

Let's see if we indeed have differences across the hybrid index with our predicted weight loss.


### Install the package

```
##
## * checking for file '/tmp/Rtmp7QcQEB/remotes3cf301188cf1a5/alicebalard-parasiteLoad-1b43216/DESCRIPTION' ... OK
## * preparing 'parasiteLoad':
## * checking DESCRIPTION meta-information ... OK
## * checking for LF line-endings in source and make files and shell scripts
## * checking for empty or unneeded directories
## * building 'parasiteLoad_0.1.0.tar.gz'
```

### Data diagnostics

#### Visualizations

```
Field %>% ggplot(aes(x = predictions_field)) +
  geom_histogram(binwidth = 1.5)
```




What is the distribution of the predicted weight loss?

Rough graph of our predictions against the hybrid index and against the

```
Field %>%
  ggplot(aes(x = HI , y = predictions_field , color = Sex)) +
  geom_smooth() +
  geom_point()
```


body length

```
## `geom_smooth()` using method = 'loess' and formula = 'y ~ x'
```



```
Field %>%
  ggplot(aes(x = Body_Length , y = predictions_field , color = Sex)) +
  geom_smooth() +
  geom_point()

## `geom_smooth()` using method = 'loess' and formula = 'y ~ x'
## Warning: Removed 1 rows containing non-finite values (`stat_smooth()`).
## Warning: Removed 1 rows containing missing values (`geom_point()`).
```



### Fitting distributions??

Ratios / Percentages are not normally distributed. Weibull is a good distributions.


Alice used weibull for the qpcr data. (paper)

```
Field <- Field %>%
  dplyr::mutate(WL = predictions_field)

x <- Field$WL

descdist(data = x, discrete = FALSE)
```


Cullen and Frey graph



```
## summary statistics
## -----
## min: 3.959833 max: 20.1775
## median: 10.09201
## mean: 10.3217
## estimated sd: 2.516379
## estimated skewness: 0.4392375
## estimated kurtosis: 3.497443
```

```
descdist(data = x, discrete = FALSE, #data is continuous
          boot = 1000)
```

Cullen and Frey graph



```
## summary statistics
## -----
## min: 3.959833 max: 20.1775
## median: 10.09201
## mean: 10.3217
## estimated sd: 2.516379
## estimated skewness: 0.4392375
## estimated kurtosis: 3.497443
```

### Test for binomial distribution

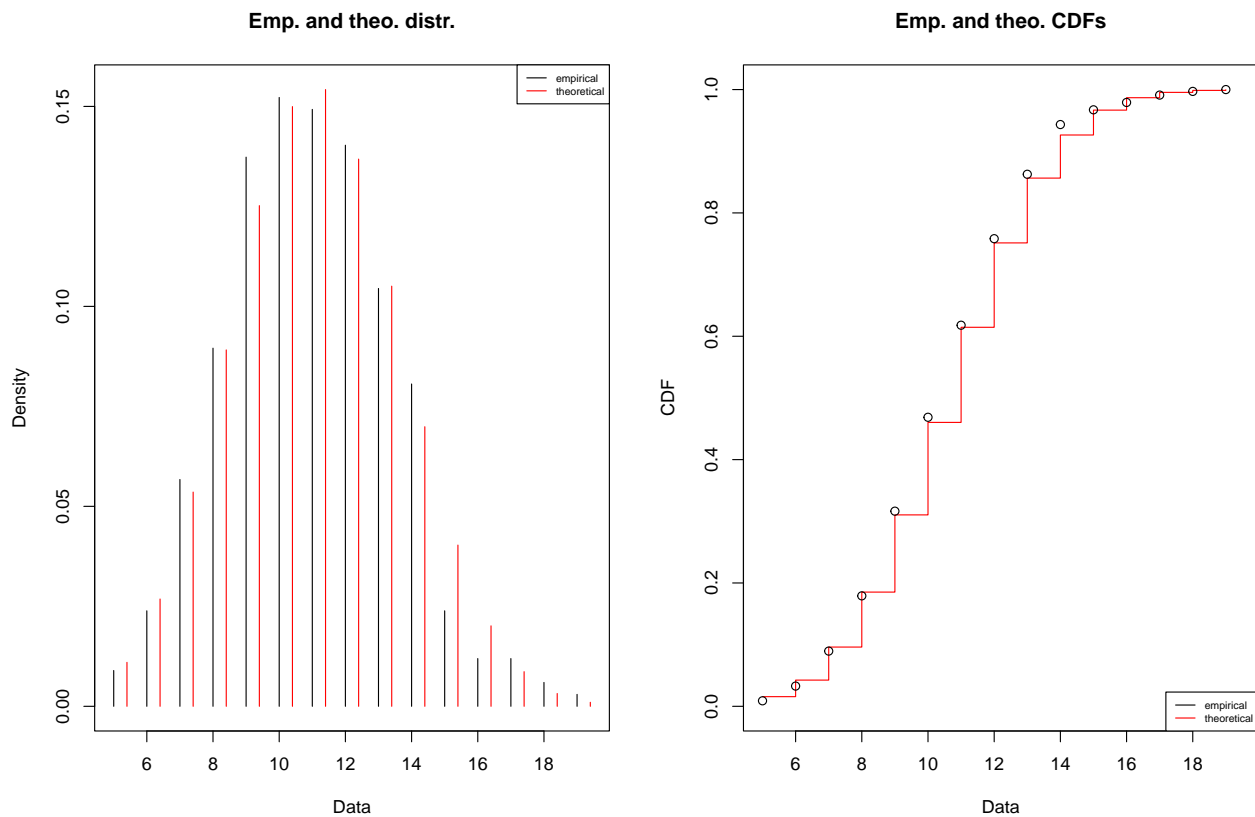
```
set.seed(10)
n = 25
size = 27
prob = .4
data = rbinom(x, size = size, prob = prob)
fit = fitdist(data = data, dist="binom",
              fix.arg=list(size = size),
              start=list(prob = 0.1))

summary(fit)
```

```
## Fitting of the distribution ' binom ' by maximum likelihood
## Parameters :
##      estimate Std. Error
## prob 0.399558 0.005150141
## Fixed parameters:
##      value
```



```
## size      27
## Loglikelihood: -779.317   AIC: 1560.634   BIC: 1564.448
plot(fit)
```



```
normal_ <- fitdist(x, "norm")
weibull_ <- fitdist(x, "weibull")
gamma_ <- fitdist(x, "gamma")

# Define function to be used to test, get the log lik and aic
tryDistrib <- function(x, distrib){
  # deals with fitdistr error:
  fit <- tryCatch(MASS::fitdistr(x, distrib), error=function(err) "fit failed")
  return(list(fit = fit,
              loglik = tryCatch(fit$loglik, error=function(err) "no loglik computed"),
              AIC = tryCatch(fit$aic, error=function(err) "no aic computed")))
}

findGoodDist <- function(x, distribs, distribs2){
  l =lapply(distribs, function(i) tryDistrib(x, i))
  names(l) <- distribs
  print(l)
  listDist <- lapply(distribs2, function(i){
    if (i %in% "t"){
```

```

    fitdistrplus::fitdist(x, i, start = list(df =2))
  } else {
    fitdistrplus::fitdist(x,i)
  }}
)
par(mfrow=c(2,2))
denscomp(listDistr, legendtext=distrib2)
cdfcomp(listDistr, legendtext=distrib2)
qqcomp(listDistr, legendtext=distrib2)
ppcomp(listDistr, legendtext=distrib2)
par(mfrow=c(1,1))
}

```

```
tryDistrib(x, "normal")
```

### Functions for testing distributions

```

## $fit
##      mean      sd
##  10.32170304  2.51261999
##  ( 0.13727909) ( 0.09707098)
##
## $loglik
## [1] -783.9886
##
## $AIC
## NULL

```

```
tryDistrib(x, "binomial")
```

```

## $fit
## [1] "fit failed"
##
## $loglik
## [1] "no loglik computed"
##
## $AIC
## [1] "no aic computed"

```

```
tryDistrib(x, "student")
```

```

## $fit
## [1] "fit failed"
##
## $loglik
## [1] "no loglik computed"
##
## $AIC
## [1] "no aic computed"

```

```
tryDistrib(x, "weibull")
```

```
## Warning in densfun(x, parm[1], parm[2], ...): NaNs produced
```

```
## Warning in densfun(x, parm[1], parm[2], ...): NaNs produced
```





```
## $fit
##      shape      scale
##  4.3023202  11.3019120
## ( 0.1703402) ( 0.1519404)
##
## $loglik
## [1] -792.4037
##
## $AIC
## NULL
```

```
tryDistrib(x, "weibullshifted")
```





```
## $fit
## [1] "fit failed"
##
## $loglik
## [1] "no loglik computed"
##
## $AIC
## [1] "no aic computed"
```

```
findGoodDist(x, "normal", "weibull")
```

```
## $normal
## $normal$fit
##      mean      sd
##  10.32170304  2.51261999
## ( 0.13727909) ( 0.09707098)
##
## $normal$loglik
## [1] -783.9886
##
## $normal$AIC
## NULL
```



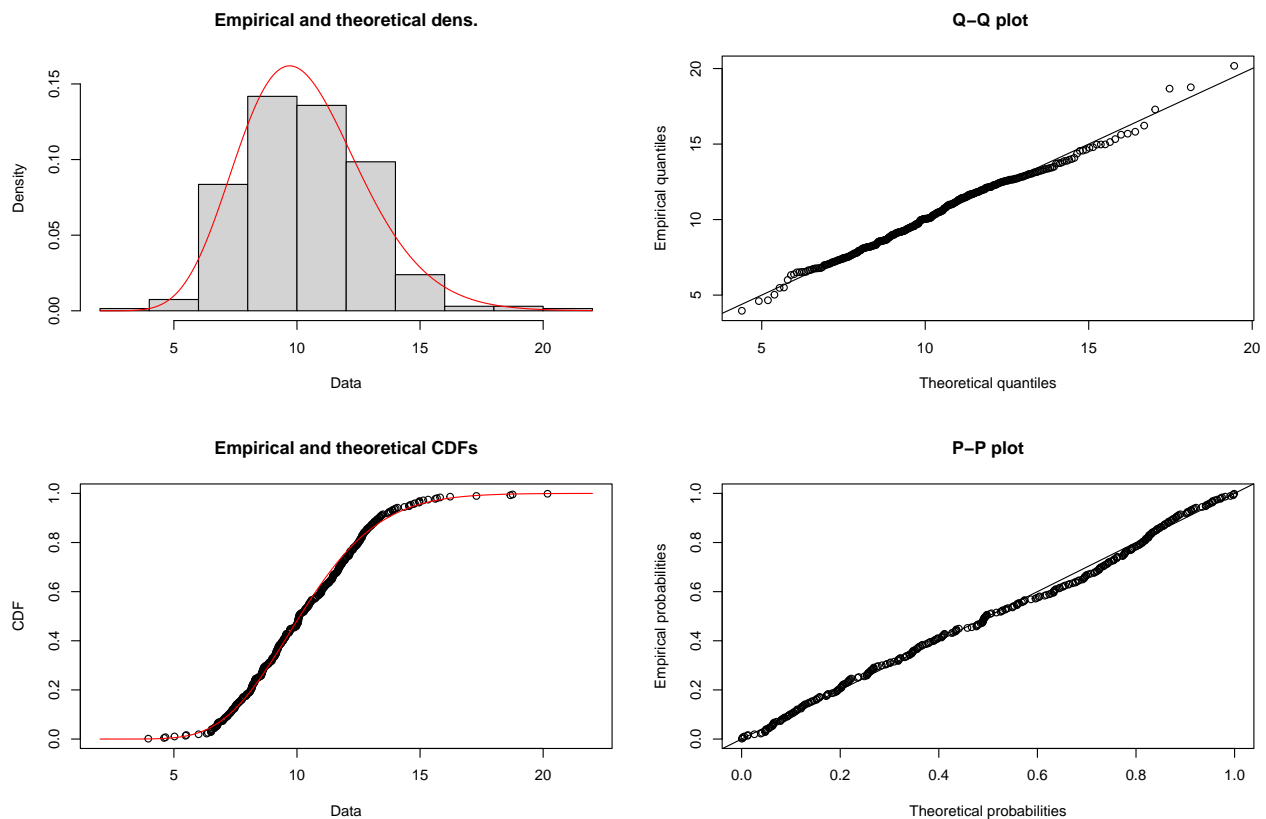
`plot(normal_)`



```
summary(normal_)
```

```
## Fitting of the distribution ' norm ' by maximum likelihood
## Parameters :
##      estimate Std. Error
## mean 10.32170 0.13727909
## sd    2.51262 0.09707091
## Loglikelihood: -783.9886   AIC: 1571.977   BIC: 1579.606
## Correlation matrix:
##      mean sd
## mean  1  0
## sd    0  1
```

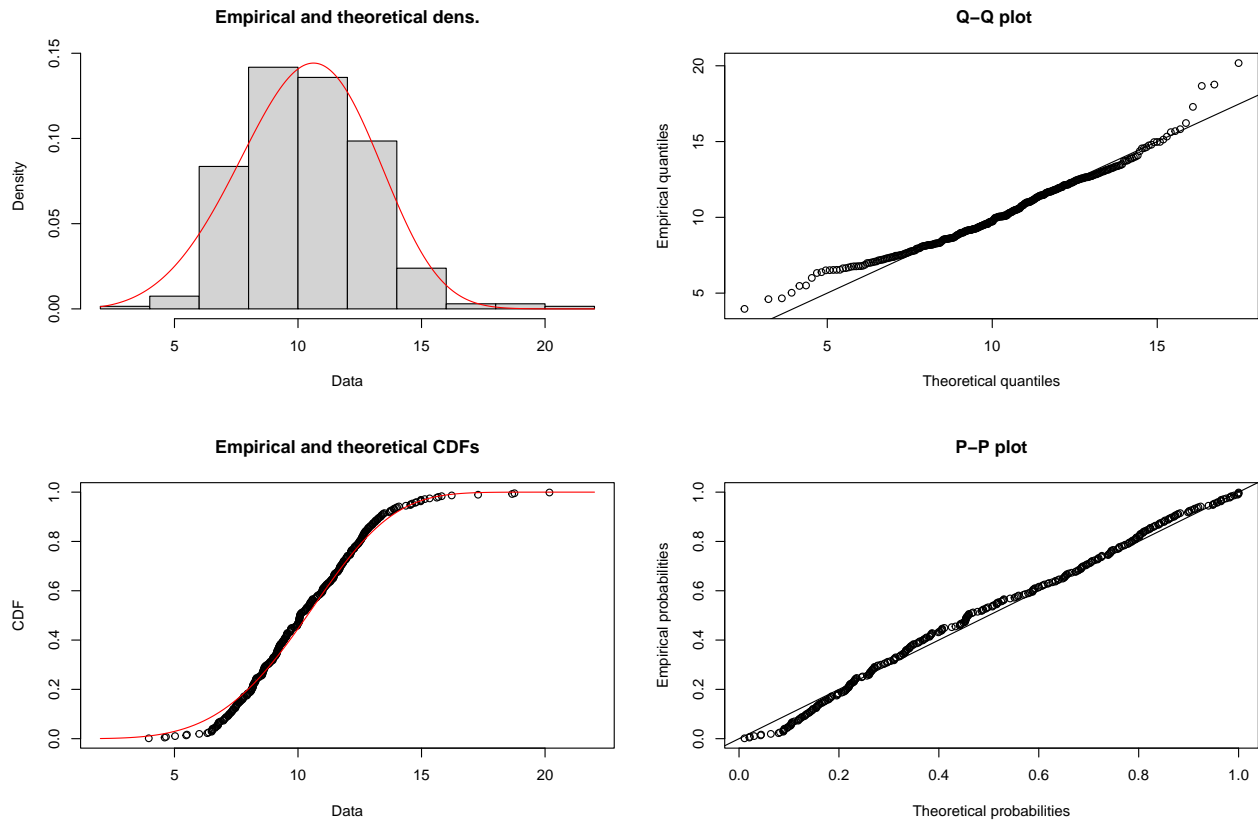
```
plot(gamma_)
```



```
summary(gamma_)
```

```
## Fitting of the distribution ' gamma ' by maximum likelihood
## Parameters :
##      estimate Std. Error
## shape 16.679232 1.2760602
## rate  1.615965 0.1255067
## Loglikelihood: -779.1252   AIC: 1562.25   BIC: 1569.879
## Correlation matrix:
##      shape      rate
## shape 1.0000000 0.9850536
## rate  0.9850536 1.0000000
```

```
plot(weibull_)
```



```
summary(weibull_)
```

```
## Fitting of the distribution ' weibull ' by maximum likelihood
## Parameters :
##      estimate Std. Error
## shape  4.302738  0.1703492
## scale 11.301733  0.1519233
## Loglikelihood: -792.4038   AIC: 1588.808   BIC: 1596.436
## Correlation matrix:
##      shape      scale
## shape 1.0000000 0.3283846
## scale 0.3283846 1.0000000
```

Is alpha significant for each hypothesis?

```
Field$Sex <- as.factor(Field$Sex)
```

```
parasiteLoad::getParamBounds("normal", data = Field, response = "WL")
```

```
##      L1start      L1LB      L1UB      L2start      L2LB      L2UB
## 10.321703036  3.959832716 20.177504703 10.321703036  3.959832716 20.177504703
##      alphaStart      alphaLB      alphaUB      mysdStart      mysdLB      mysdUB
##      0.000000000 -5.000000000  5.000000000  1.000000000  0.000000001 10.000000000
```

```

speparam <- c(L1start = 10,
              L1LB = 1e-9,
              L1UB = 20,
              L2start = 10,
              L2LB = 1e-9,
              L2UB = 20,
              alphaStart = 0, alphaLB = -5, alphaUB = 5,
              myshapeStart = 1, myshapeLB = 1e-9, myshapeUB = 5)

```

```
##All
```

```

fitWL_Sex <- parasiteLoad::analyse(data = Field,
                                   response = "WL",
                                   model = "normal",
                                   group = "Sex")

```

```

## [1] "Analysing data for response: WL"
## [1] "Fit for the response: WL"
## [1] "Fitting for all"
## [1] "Fitting model basic without alpha"
## [1] "Did converge"
## [1] "Fitting model basic with alpha"
## [1] "Did converge"
## [1] "Fitting model advanced without alpha"
## [1] "Did converge"
## [1] "Fitting model advanced with alpha"
## [1] "Did converge"
## [1] "Fitting for groupA : F"
## [1] "Fitting model basic without alpha"
## [1] "Did converge"
## [1] "Fitting model basic with alpha"
## [1] "Did converge"
## [1] "Fitting model advanced without alpha"
## [1] "Did converge"
## [1] "Fitting model advanced with alpha"
## [1] "Did converge"
## [1] "Fitting for groupB : M"
## [1] "Fitting model basic without alpha"
## [1] "Did converge"
## [1] "Fitting model basic with alpha"
## [1] "Did converge"
## [1] "Fitting model advanced without alpha"
## [1] "Did converge"
## [1] "Fitting model advanced with alpha"
## [1] "Did converge"
## [1] "Testing H0 no alpha vs alpha"
##      dLL dDF      pvalue
## 1 2.63    1 0.02185549
## [1] "Testing H1 no alpha vs alpha"
##      dLL dDF      pvalue
## 1 2.01    1 0.04514727
## [1] "Testing H2 groupA no alpha vs alpha"
##      dLL dDF      pvalue
## 1 0.99    1 0.1592538
## [1] "Testing H2 groupB no alpha vs alpha"

```

```
##      dLL dDF      pvalue
## 1 1.81    1 0.05725558
## [1] "Testing H3 groupA no alpha vs alpha"
##      dLL dDF      pvalue
## 1 1.26    1 0.1130798
## [1] "Testing H3 groupB no alpha vs alpha"
##      dLL dDF      pvalue
## 1 1.82    1 0.05636448
## [1] "Testing H1 vs H0"
##      dLL dDF      pvalue
## 1 0.92    1 0.1756223
## [1] "Testing H2 vs H0"
##      dLL dDF      pvalue
## 1 2.09    3 0.2417575
## [1] "Testing H3 vs H1"
##      dLL dDF      pvalue
## 1 5.75    4 0.02147942
## [1] "Testing H3 vs H2"
##      dLL dDF      pvalue
## 1 4.57    2 0.01032858
```

```
parasiteLoad::analyse(data = Field,
                      response = "WL",
                      model = "normal",
                      group = "Sex")
```

```
## [1] "Analysing data for response: WL"
## [1] "Fit for the response: WL"
## [1] "Fitting for all"
## [1] "Fitting model basic without alpha"
## [1] "Did converge"
## [1] "Fitting model basic with alpha"
## [1] "Did converge"
## [1] "Fitting model advanced without alpha"
## [1] "Did converge"
## [1] "Fitting model advanced with alpha"
## [1] "Did converge"
## [1] "Fitting for groupA : F"
## [1] "Fitting model basic without alpha"
## [1] "Did converge"
## [1] "Fitting model basic with alpha"
## [1] "Did converge"
## [1] "Fitting model advanced without alpha"
## [1] "Did converge"
## [1] "Fitting model advanced with alpha"
## [1] "Did converge"
## [1] "Fitting for groupB : M"
## [1] "Fitting model basic without alpha"
## [1] "Did converge"
## [1] "Fitting model basic with alpha"
## [1] "Did converge"
## [1] "Fitting model advanced without alpha"
## [1] "Did converge"
## [1] "Fitting model advanced with alpha"
## [1] "Did converge"
```



```

## [1] "Testing H0 no alpha vs alpha"
##      dLL dDF      pvalue
## 1 2.63    1 0.02185549
## [1] "Testing H1 no alpha vs alpha"
##      dLL dDF      pvalue
## 1 2.01    1 0.04514727
## [1] "Testing H2 groupA no alpha vs alpha"
##      dLL dDF      pvalue
## 1 0.99    1 0.1592538
## [1] "Testing H2 groupB no alpha vs alpha"
##      dLL dDF      pvalue
## 1 1.81    1 0.05725558
## [1] "Testing H3 groupA no alpha vs alpha"
##      dLL dDF      pvalue
## 1 1.26    1 0.1130798
## [1] "Testing H3 groupB no alpha vs alpha"
##      dLL dDF      pvalue
## 1 1.82    1 0.05636448
## [1] "Testing H1 vs H0"
##      dLL dDF      pvalue
## 1 0.92    1 0.1756223
## [1] "Testing H2 vs H0"
##      dLL dDF      pvalue
## 1 2.09    3 0.2417575
## [1] "Testing H3 vs H1"
##      dLL dDF      pvalue
## 1 5.75    4 0.02147942
## [1] "Testing H3 vs H2"
##      dLL dDF      pvalue
## 1 4.57    2 0.01032858

## $H0
##
## Call:
## bbmle::mle2(minuslogl = response ~ dnorm(mean = MeanLoad(L1,
##      L1, alpha, HI), sd = mysd), start = start, method = config$method,
##      optimizer = config$optimizer, data = data, lower = c(L1 = paramBounds[["L1LB"]],
##      mysd = paramBounds[["mysdLB"]], alpha = paramBounds[["alphaLB"]]),
##      upper = c(L1 = paramBounds[["L1UB"]], mysd = paramBounds[["mysdUB"]],
##      alpha = paramBounds[["alphaUB"]]), control = config$control)
##
## Coefficients:
##      L1      mysd      alpha
## 9.7996797 2.4929816 -0.2305268
##
## Log-likelihood: -781.36
## Best method: bobyqa
##
## $H1
##
## Call:
## bbmle::mle2(minuslogl = response ~ dnorm(mean = MeanLoad(L1,
##      L2, alpha, HI), sd = mysd), start = start, method = config$method,
##      optimizer = config$optimizer, data = data, lower = c(L1 = paramBounds[["L1LB"]],

```

```

##      mysd = paramBounds[["mysdLB"]], L2 = paramBounds[["L2LB"]],
##      alpha = paramBounds[["alphaLB"]]), upper = c(L1 = paramBounds[["L1UB"]],
##      mysd = paramBounds[["mysdUB"]], L2 = paramBounds[["L2UB"]],
##      alpha = paramBounds[["alphaUB"]]), control = config$control)
##
## Coefficients:
##      L1      L2      alpha      mysd
##  9.564795 10.071590 -0.204059  2.486166
##
## Log-likelihood: -780.44
## Best method: bobyqa
##
## $H2
## $H2$groupA
##
## Call:
## bbmle::mle2(minuslogl = response ~ dnorm(mean = MeanLoad(L1,
##      L1, alpha, HI), sd = mysd), start = start, method = config$method,
##      optimizer = config$optimizer, data = data, lower = c(L1 = paramBounds[["L1LB"]],
##      mysd = paramBounds[["mysdLB"]], alpha = paramBounds[["alphaLB"]]),
##      upper = c(L1 = paramBounds[["L1UB"]], mysd = paramBounds[["mysdUB"]],
##      alpha = paramBounds[["alphaUB"]]), control = config$control)
##
## Coefficients:
##      L1      mysd      alpha
## 10.0350054  2.6318415 -0.1997412
##
## Log-likelihood: -398.57
## Best method: bobyqa
##
## $H2$groupB
##
## Call:
## bbmle::mle2(minuslogl = response ~ dnorm(mean = MeanLoad(L1,
##      L1, alpha, HI), sd = mysd), start = start, method = config$method,
##      optimizer = config$optimizer, data = data, lower = c(L1 = paramBounds[["L1LB"]],
##      mysd = paramBounds[["mysdLB"]], alpha = paramBounds[["alphaLB"]]),
##      upper = c(L1 = paramBounds[["L1UB"]], mysd = paramBounds[["mysdUB"]],
##      alpha = paramBounds[["alphaUB"]]), control = config$control)
##
## Coefficients:
##      L1      mysd      alpha
##  9.5523403  2.3329417 -0.2688342
##
## Log-likelihood: -380.7
## Best method: bobyqa
##
## $H3
## $H3$groupA
##
## Call:
## bbmle::mle2(minuslogl = response ~ dnorm(mean = MeanLoad(L1,
##      L2, alpha, HI), sd = mysd), start = start, method = config$method,

```

```

##      optimizer = config$optimizer, data = data, lower = c(L1 = paramBounds[["L1LB"]],
##      mysd = paramBounds[["mysdLB"]], L2 = paramBounds[["L2LB"]],
##      alpha = paramBounds[["alphaLB"]]), upper = c(L1 = paramBounds[["L1UB"]],
##      mysd = paramBounds[["mysdUB"]], L2 = paramBounds[["L2UB"]],
##      alpha = paramBounds[["alphaUB"]]), control = config$control)
##
## Coefficients:
##      L1      L2      alpha      mysd
## 10.1934433  9.7543876 -0.2423411  2.6273997
##
## Log-likelihood: -398.28
## Best method: bobyqa
##
## $H3$groupB
##
## Call:
## bbmle::mle2(minuslogl = response ~ dnorm(mean = MeanLoad(L1,
##      L2, alpha, HI), sd = mysd), start = start, method = config$method,
##      optimizer = config$optimizer, data = data, lower = c(L1 = paramBounds[["L1LB"]],
##      mysd = paramBounds[["mysdLB"]], L2 = paramBounds[["L2LB"]],
##      alpha = paramBounds[["alphaLB"]]), upper = c(L1 = paramBounds[["L1UB"]],
##      mysd = paramBounds[["mysdUB"]], L2 = paramBounds[["L2UB"]],
##      alpha = paramBounds[["alphaUB"]]), control = config$control)
##
## Coefficients:
##      L1      L2      alpha      mysd
##  8.7485937 10.1460223 -0.2653402  2.2741121
##
## Log-likelihood: -376.41
## Best method: bobyqa
plot_WL_Sex<- bananaPlot(mod = fitWL_Sex$H3,
      data = Field,
      response = "WL",
      group = "Sex") +
      scale_fill_manual(values = c("blueviolet", "limegreen")) +
      scale_color_manual(values = c("blueviolet", "limegreen")) +
      theme_bw()

## Scale for fill is already present.
## Adding another scale for fill, which will replace the existing scale.
## Scale for colour is already present.
## Adding another scale for colour, which will replace the existing scale.

# Create HI bar
HIgradientBar <- ggplot(data.frame(hi = seq(0,1,0.0001)),
      aes(x=hi, y=1, fill = hi)) +
      geom_tile() +
      theme_void() +
      scale_fill_gradient(low = "blue", high = "red") +
      scale_x_continuous(expand=c(.01,0)) +
      scale_y_continuous(expand=c(0,0)) +
      theme(legend.position = 'none')


plot_grid(plot_WL_Sex,

```


```

HIgradientBar,
nrow = 2,
align = "v",
axis = "tlr",
rel_heights = c(13, 1))

```



plot\_WL\_Sex



H0: the expected load for the subspecies and between 2 groups is the same

H1: the mean load across 2 groups is the same, but can differ across subspecies

H2: the mean load across subspecies is the same, but can differ between the 2 groups


H3: the mean load can differ both across subspecies and between 2 groups

```
ggplot(data = Field, aes(x = delta_ct_cewe_MminusE, y = WL)) +
  geom_point() +
  stat_smooth(method= "lm")
```

```
## `geom_smooth()` using formula = 'y ~ x'
```

```
## Warning: Removed 150 rows containing non-finite values (`stat_smooth()`).
```

```
## Warning: Removed 150 rows containing missing values (`geom_point()`).
```



```
Field2 <- Field %>%
  drop_na(delta_ct_cewe_MminusE)

cor(Field2$WL, Field2$delta_ct_cewe_MminusE)

## [1] 0.1576532

tolerance <- lm(WL ~ delta_ct_cewe_MminusE, data = Field)

summary(tolerance)

##
## Call:
## lm(formula = WL ~ delta_ct_cewe_MminusE, data = Field)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -5.2762 -1.8206 -0.3349  1.6904  8.7213
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    10.91325    0.45066   24.22  <2e-16 ***
## delta_ct_cewe_MminusE  0.11288    0.05227    2.16  0.0321 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.341 on 183 degrees of freedom
```


```
## (150 observations deleted due to missingness)
## Multiple R-squared:  0.02485,    Adjusted R-squared:  0.01953
## F-statistic: 4.664 on 1 and 183 DF,  p-value: 0.0321
```

```
confint(tolerance)
```

```
##                2.5 %    97.5 %
## (Intercept)      10.024092567 11.802405
## delta_ct_cewe_MminusE 0.009757681 0.216005
```

```
ggplot(data = Field, aes(x = OPG, y = WL)) +
  geom_point() +
  stat_smooth(method= "lm") +
  scale_x_log10()
```

```
## Warning: Transformation introduced infinite values in continuous x-axis
## Transformation introduced infinite values in continuous x-axis
## `geom_smooth()` using formula = 'y ~ x'
## Warning: Removed 280 rows containing non-finite values (`stat_smooth()`).
## Warning: Removed 157 rows containing missing values (`geom_point()`).
```



```
Field2 <- Field %>%
  drop_na(OPG)

cor(Field2$WL, Field2$OPG)
```

```
## [1] 0.08025742
```

```
tolerance <- lm(WL ~ OPG, data = Field)
```

```
summary(tolerance)
```

```
##
## Call:
## lm(formula = WL ~ OPG, data = Field)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -6.5452 -1.9797 -0.0749  1.7505  9.6588
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 1.051e+01  2.012e-01  52.225  <2e-16 ***
## OPG          3.732e-08  3.493e-08   1.068    0.287
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.65 on 176 degrees of freedom
## (157 observations deleted due to missingness)
## Multiple R-squared:  0.006441, Adjusted R-squared:  0.000796
## F-statistic: 1.141 on 1 and 176 DF, p-value: 0.2869
```

```
confint(tolerance)
```

```
##              2.5 %      97.5 %
## (Intercept) 1.010806e+01 1.090201e+01
## OPG          -3.162786e-08 1.062601e-07
```

```
tolerance <- lm(WL ~ OPG * delta_ct_cewe_MminusE, data = Field)
```

```
summary(tolerance)
```

```
##
## Call:
## lm(formula = WL ~ OPG * delta_ct_cewe_MminusE, data = Field)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4.6452 -1.9629 -0.4451  1.4172  8.6849
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      1.146e+01  9.024e-01  12.699  <2e-16 ***
## OPG              -1.183e-05  2.335e-05  -0.506    0.615
## delta_ct_cewe_MminusE  1.787e-01  1.181e-01   1.512    0.137
## OPG:delta_ct_cewe_MminusE  5.325e-06  6.842e-06   0.778    0.440
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.571 on 45 degrees of freedom
## (286 observations deleted due to missingness)
```



```
## Multiple R-squared:  0.07703,    Adjusted R-squared:  0.0155
## F-statistic: 1.252 on 3 and 45 DF,  p-value: 0.3023
```

```
confint(tolerance)
```

```
##                2.5 %      97.5 %
## (Intercept)      9.641497e+00 1.327639e+01
## OPG              -5.885891e-05 3.520486e-05
## delta_ct_cewe_MminusE -5.927564e-02 4.166351e-01
## OPG:delta_ct_cewe_MminusE -8.455062e-06 1.910478e-05
```


```
Field <- Field %>%
  dplyr::mutate(BMI = Body_Weight / Body_Length)
```

```
ggplot(data = Field, aes(x = BMI, y = WL)) +
  geom_point() +
  stat_smooth(method= "lm")
```

```
## `geom_smooth()` using formula = 'y ~ x'
```

```
## Warning: Removed 1 rows containing non-finite values (`stat_smooth()`).
```

```
## Warning: Removed 1 rows containing missing values (`geom_point()`).
```



```
bmi <- lm(WL ~ BMI, data = Field)
```

```
cor(Field$BMI, Field$WL, use = "complete.obs")
```

```
## [1] -0.1138541
```

```
summary(bmi)
```

```
##
## Call:
## lm(formula = WL ~ BMI, data = Field)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -5.8747 -1.8728 -0.0796  1.7947 10.0742
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  11.6905     0.6707   17.431  <2e-16 ***
## BMI          -7.3938     3.5409   -2.088   0.0376 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.507 on 332 degrees of freedom
## (1 observation deleted due to missingness)
## Multiple R-squared:  0.01296,    Adjusted R-squared:  0.00999
## F-statistic:  4.36 on 1 and 332 DF,  p-value: 0.03755
```

```
confint(bmi)
```

```
##              2.5 %      97.5 %
## (Intercept)  10.37122 13.0098096
## BMI          -14.35928 -0.4283224
```