

## 10. Applying random forest on field data - gene

Fay

2022-11-04

### Aim:

- Applying the models established in the script: 9
- How are hybrid mice different to the parental species?

### Load necessary libraries:

```
#install.packages("optimx", version = "2021-10.12") # this package is required for  
#the parasite load package to work  
library(tidyverse)  
library(tidyr)  
library(dplyr)  
library(cowplot)  
library(randomForest)  
library(ggplot2)  
library(VIM) # visualizing missing data  
library(mice) # imputing missing data without predictors  
library(ggpubr)  
library(optimx)  
library(rfUtilities) # Implements a permutation test cross-validation for  
# Random Forests models  
library(mice) #imputations  
library(fitdistrplus) #testing distributions  
library(logspline)  
library(caret)
```

### Field data

#### Import field data

```
hm <- read.csv("output_data/imputed_mice.csv")
```

#### Clean data

```
Field <- hm %>%  
  filter(origin == "Field") %>%  
  drop_na(HI)
```

We have 1921 mice in total.

```
EqPCR.cols      <- c("delta_ct_cewe_MminusE", "MC.Eimeria", "Ct.Eimeria") #, "Ct.Mus" "delta_ct_ilwe_Mmin
Genes_wild      <- c("IFNy", "CXCR3", "IL.6", "IL.13", "IL.10",
                    "IL1RN", "CASP1", "CXCL9", "IDO1", "IRGM1", "MPO",
                    "MUC2", "MUC5AC", "MYD88", "NCR1", "PRF1", "RETNLB", "SOCS1",
                    "TICAM1", "TNF") #, "IL.12", "IRG6")
```

Prepare vectors for selecting

Actual Cleaning

```
#select the imputed gene columns
gene <- Field %>%
  dplyr::select(c(Mouse_ID, "IFNy", "CXCR3", "IL.6", "IL.13", "IL.10",
                  "IL1RN", "CASP1", "CXCL9", "IDO1", "IRGM1", "MPO",
                  "MUC2", "MUC5AC", "MYD88", "NCR1", "PRF1", "RETNLB", "SOCS1",
                  "TICAM1", "TNF"))

genes <- gene %>%
  dplyr::select(-Mouse_ID)

#remove rows with only nas
genes <- genes[,colSums(is.na(genes))<nrow(genes)]

#remove columns with only nas
genes <- genes[rowSums(is.na(genes)) != ncol(genes), ]

# select the same rows from the gene data
gene <- gene[row.names(genes),]

# select the same rows from the field data
Field <- Field[row.names(genes),]
```

## Predicting weight loss in our imputed field data

Start by making the predictions for the field data.

```
# load predicting weight loss model
weight_loss_predict <- readRDS("r_scripts/models/predict_WL.rds")

set.seed(540)

#The predict() function in R is used to predict the values based on the input data.
predictions_field <- predict(weight_loss_predict, genes)

#make the vector positive so that the distributions further down work
predictions_field <- predictions_field * (-1)

# assign test.data to a new object, so that we can make changes
result_field <- genes
```

```
#add the new variable of predictions to the result object
result_field <- cbind(result_field, predictions_field)

# add it to the field data
Field <- cbind(Field, predictions_field)
```

## It is time to apply the package of Alice Balard et al. on our predictions!

Let's see if we indeed have differences across the hybrid index with our predicted weight loss.

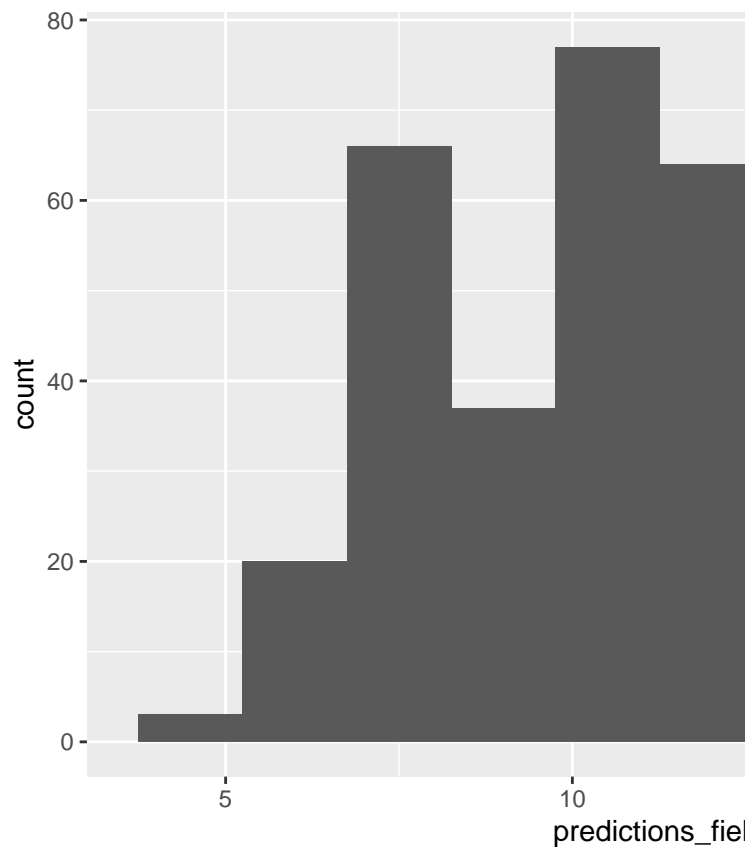
### Install the package

```
##
## * checking for file '/tmp/Rtmpu2Tx7f/remotesf18e95316a455/alicebalard-parasiteLoad-1b43216/DESCRIPTION' ... OK
## * preparing 'parasiteLoad':
## * checking DESCRIPTION meta-information ... OK
## * checking for LF line-endings in source and make files and shell scripts
## * checking for empty or unneeded directories
## * building 'parasiteLoad_0.1.0.tar.gz'
```

### Data diagnostics

#### Visualizations

```
Field %>% ggplot(aes(x = predictions_field)) +
  geom_histogram(binwidth = 1.5)
```



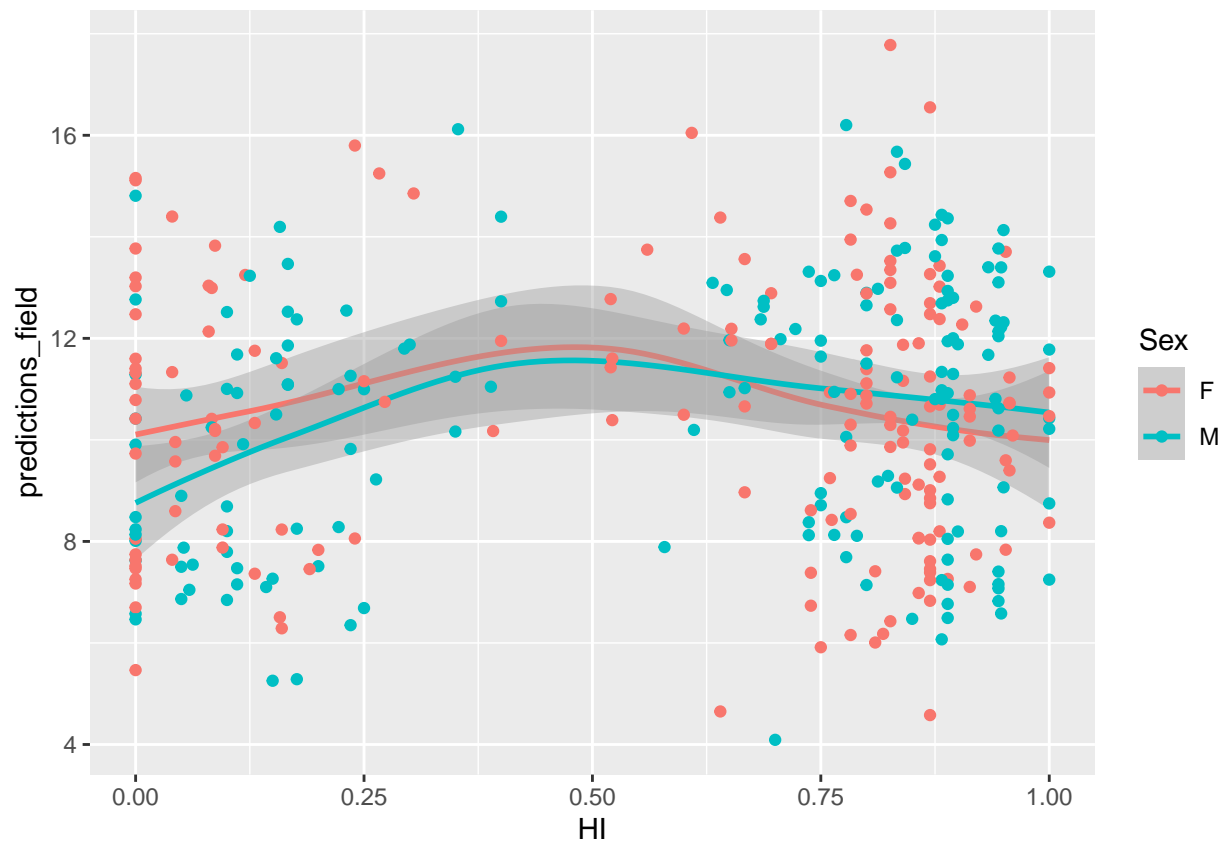
What is the distribution of the predicted weight loss?

Rough graph of our predictions against the hybrid index and against the

```
Field %>%
  ggplot(aes(x = HI , y = predictions_field , color = Sex)) +
  geom_smooth() +
  geom_point()
```

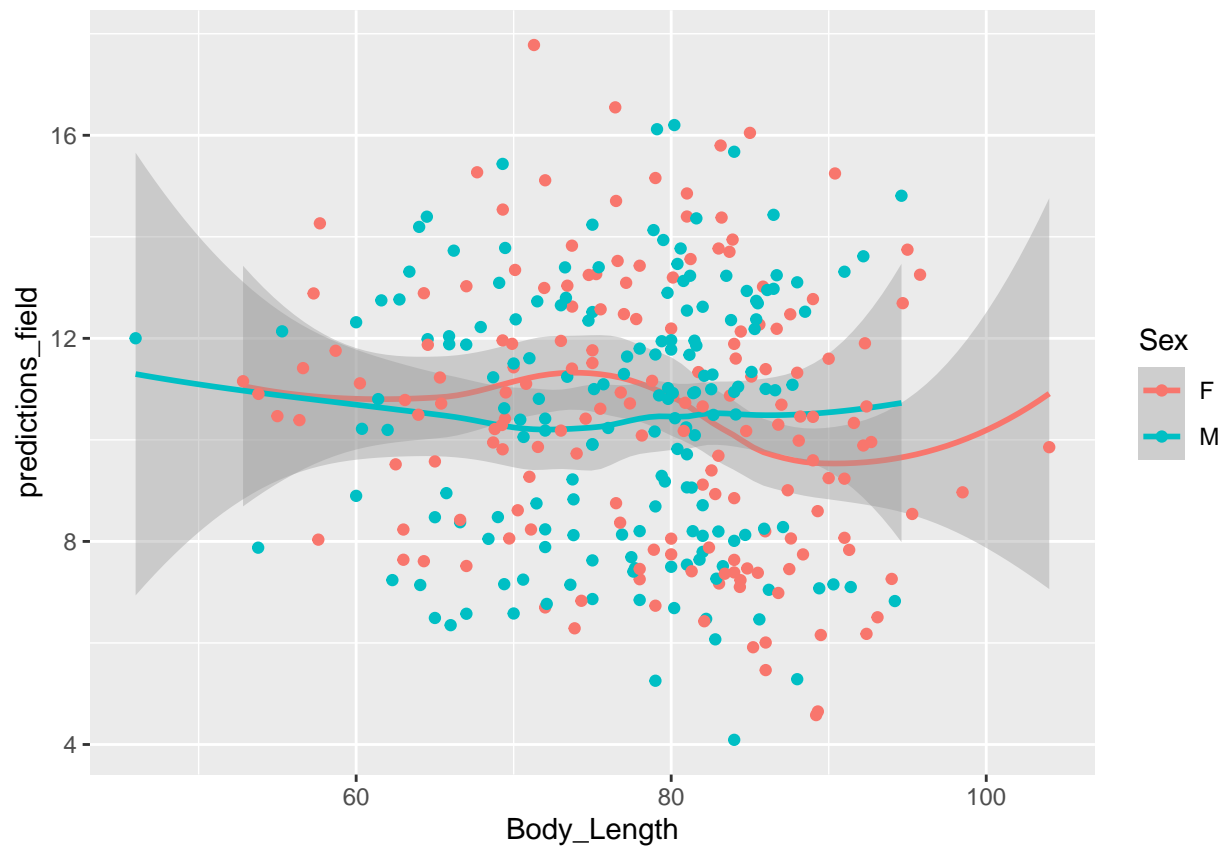
body length

```
## `geom_smooth()` using method = 'loess' and formula = 'y ~ x'
```



```
Field %>%
  ggplot(aes(x = Body_Length , y = predictions_field , color = Sex)) +
  geom_smooth() +
  geom_point()

## `geom_smooth()` using method = 'loess' and formula = 'y ~ x'
## Warning: Removed 1 rows containing non-finite values (`stat_smooth()`).
## Warning: Removed 1 rows containing missing values (`geom_point()`).
```



### Fitting distributions??

Ratios / Percentages are not normally distributed. Weibull is a good distributions.

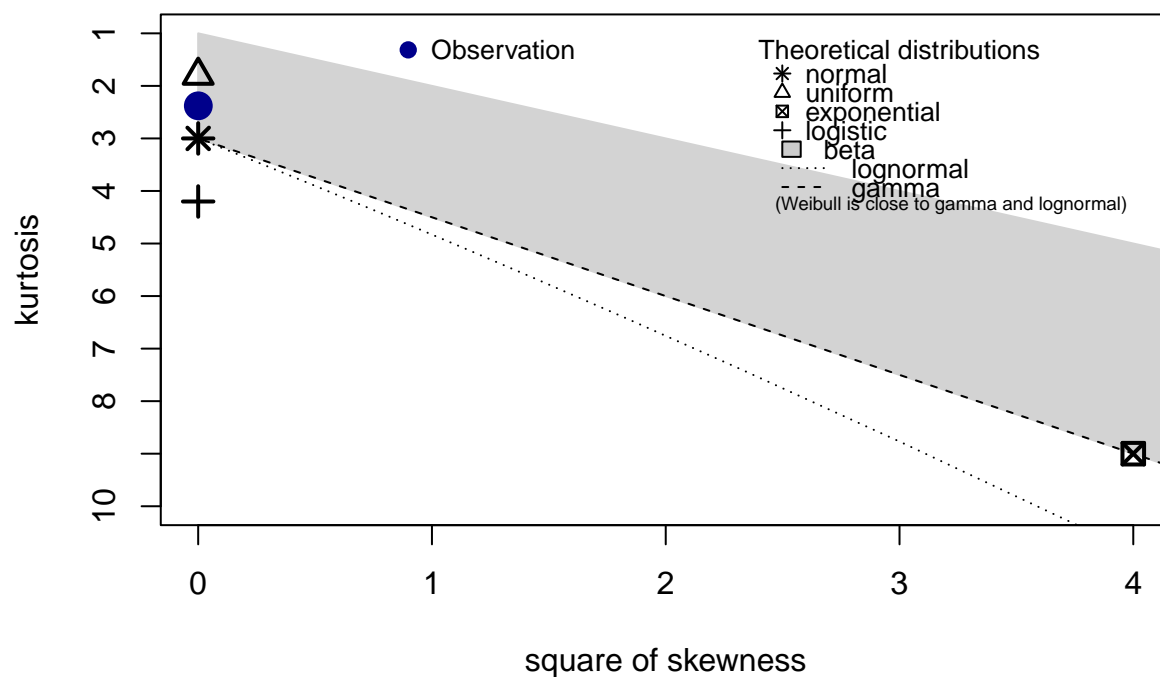
Alice used weibull for the qpcr data. (paper)

```
Field <- Field %>%
  dplyr::mutate(WL = predictions_field)

x <- Field$WL

descdist(data = x, discrete = FALSE)
```


## Cullen and Frey graph



```
## summary statistics
## -----
## min: 4.091834 max: 17.77793
## median: 10.65928
## mean: 10.4532
## estimated sd: 2.573586
## estimated skewness: 0.02495982
## estimated kurtosis: 2.379431
```

```
descdist(data = x, discrete = FALSE, #data is continuous
          boot = 1000)
```

## Cullen and Frey graph



```
## summary statistics
## -----
## min: 4.091834 max: 17.77793
## median: 10.65928
## mean: 10.4532
## estimated sd: 2.573586
## estimated skewness: 0.02495982
## estimated kurtosis: 2.379431
```

### Test for binomial distribution

```
set.seed(10)
n = 25
size = 27
prob = .4
data = rbinom(x, size = size, prob = prob)
fit = fitdist(data = data, dist="binom",
              fix.arg=list(size = size),
              start=list(prob = 0.1))

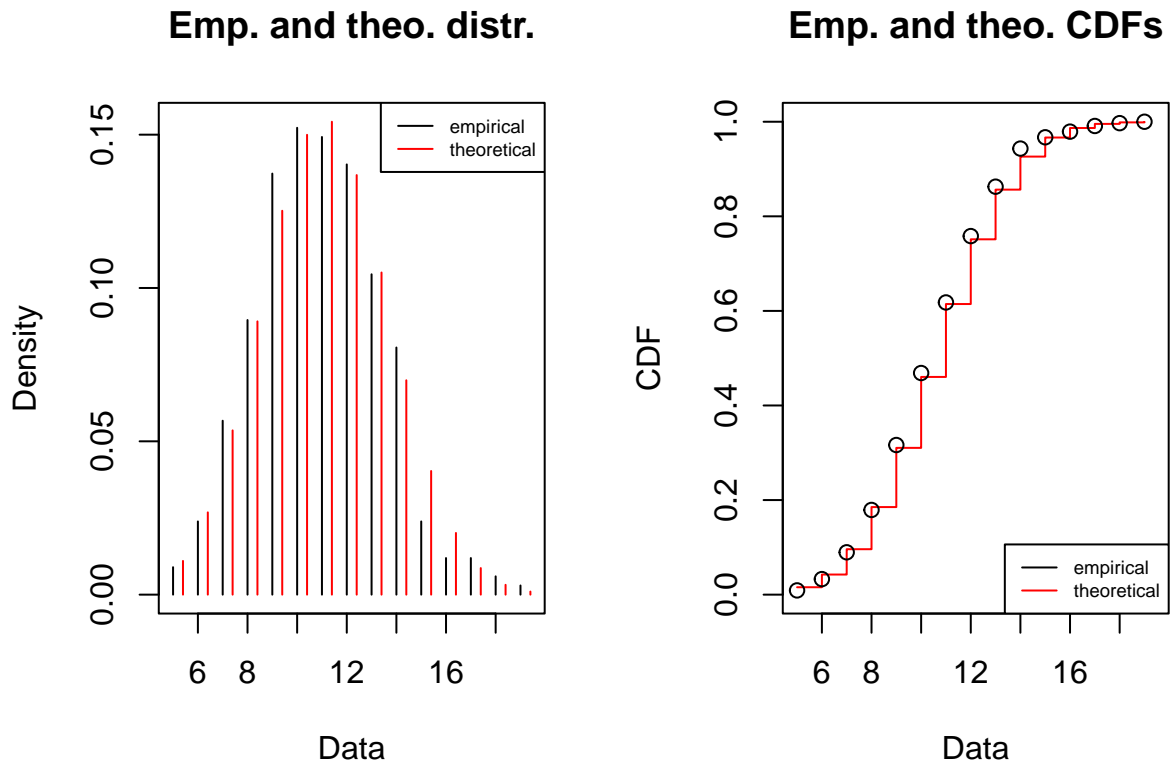
summary(fit)
```

```
## Fitting of the distribution ' binom ' by maximum likelihood
## Parameters :
##      estimate Std. Error
## prob 0.399558 0.005150141
## Fixed parameters:
##      value
## size    27
```



```
## Loglikelihood: -779.317   AIC: 1560.634   BIC: 1564.448
```

```
plot(fit)
```



```
normal_ <- fitdist(x, "norm")
weibull_ <- fitdist(x, "weibull")
gamma_ <- fitdist(x, "gamma")

# Define function to be used to test, get the log lik and aic
tryDistrib <- function(x, distrib){
  # deals with fitdistr error:
  fit <- tryCatch(MASS::fitdistr(x, distrib), error=function(err) "fit failed")
  return(list(fit = fit,
              loglik = tryCatch(fit$loglik, error=function(err) "no loglik computed"),
              AIC = tryCatch(fit$aic, error=function(err) "no aic computed")))
}

findGoodDist <- function(x, distrib, distrib2){
  l =lapply(distrib, function(i) tryDistrib(x, i))
  names(l) <- distrib
  print(l)
  listDistr <- lapply(distrib2, function(i){
    if (i %in% "t"){
      fitdistrplus::fitdist(x, i, start = list(df =2))
    } else {
      fitdistrplus::fitdist(x,i)
    }
  })
}
```

```

    }}
  )
  par(mfrow=c(2,2))
  denscomp(listDistr, legendtext=distrib2)
  cdfcomp(listDistr, legendtext=distrib2)
  qqcomp(listDistr, legendtext=distrib2)
  ppcomp(listDistr, legendtext=distrib2)
  par(mfrow=c(1,1))
}

```

```
tryDistrib(x, "normal")
```

### Functions for testing distributions

```

## $fit
##      mean      sd
## 10.45320110  2.56974179
## ( 0.14039999) ( 0.09927779)
##
## $loglik
## [1] -791.5192
##
## $AIC
## NULL

```

```
tryDistrib(x, "binomial")
```

```

## $fit
## [1] "fit failed"
##
## $loglik
## [1] "no loglik computed"
##
## $AIC
## [1] "no aic computed"

```

```
tryDistrib(x, "student")
```

```

## $fit
## [1] "fit failed"
##
## $loglik
## [1] "no loglik computed"
##
## $AIC
## [1] "no aic computed"

```

```
tryDistrib(x, "weibull")
```

```

## $fit
##      shape      scale
##  4.5546414 11.4515521
## ( 0.1927141) ( 0.1450042)
##
## $loglik
## [1] -791.1953

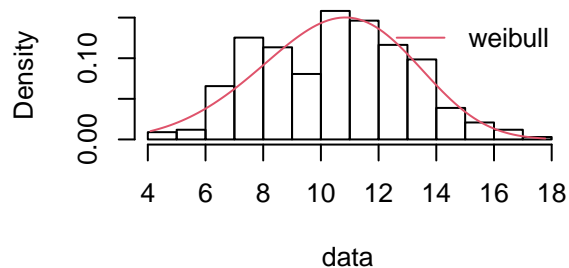
```

```
##
## $AIC
## NULL
tryDistrib(x, "weibullshifted")

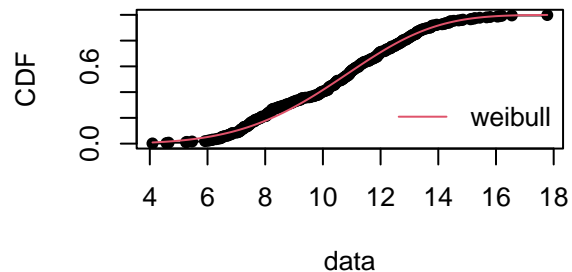
## $fit
## [1] "fit failed"
##
## $loglik
## [1] "no loglik computed"
##
## $AIC
## [1] "no aic computed"
findGoodDist(x, "normal", "weibull")

## $normal
## $normal$fit
##      mean      sd
## 10.45320110 2.56974179
## ( 0.14039999) ( 0.09927779)
##
## $normal$loglik
## [1] -791.5192
##
## $normal$AIC
## NULL
```

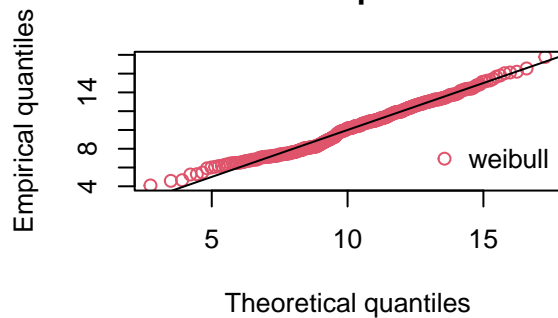
**Histogram and theoretical densities**



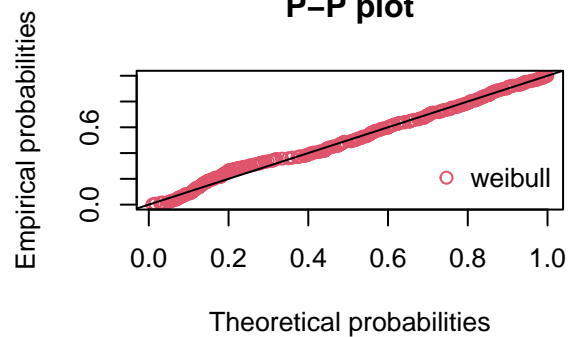
**Empirical and theoretical CDFs**



**Q-Q plot**

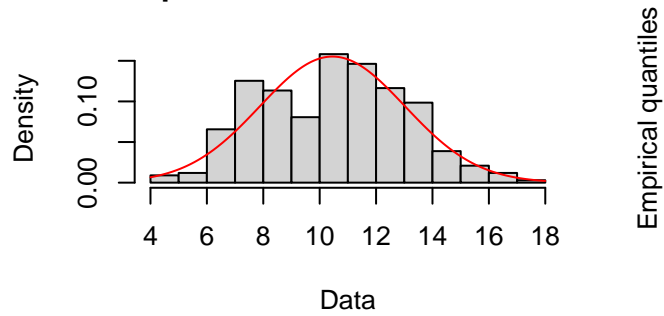


**P-P plot**

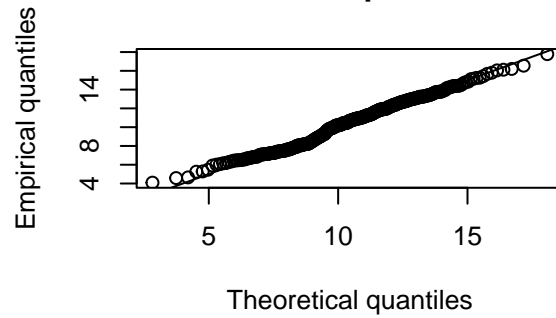


```
plot(normal_)
```

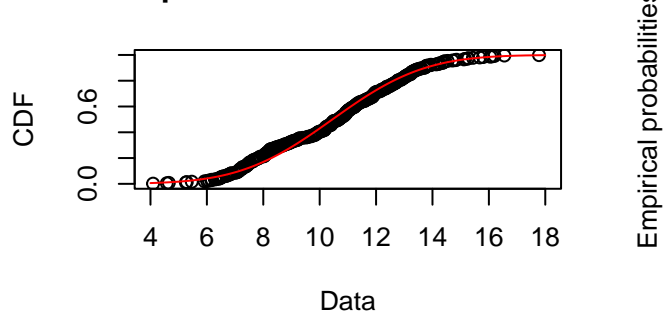
**Empirical and theoretical dens.**



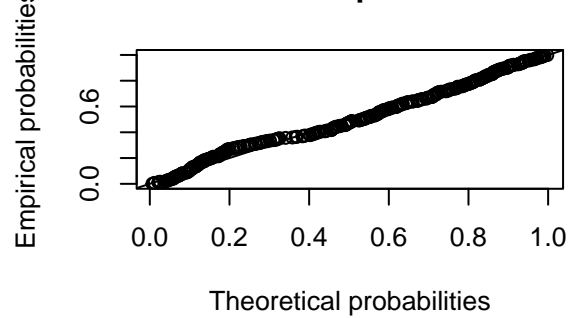
**Q-Q plot**



**Empirical and theoretical CDFs**



**P-P plot**

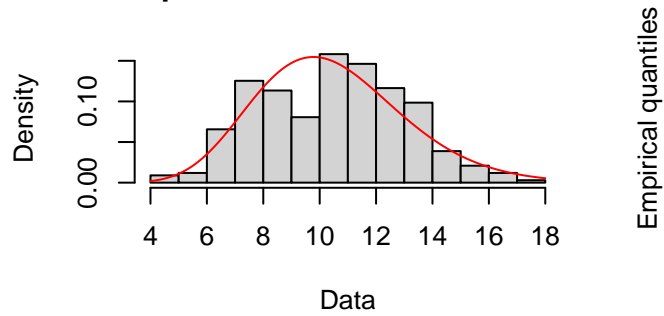


```
summary(normal_)
```

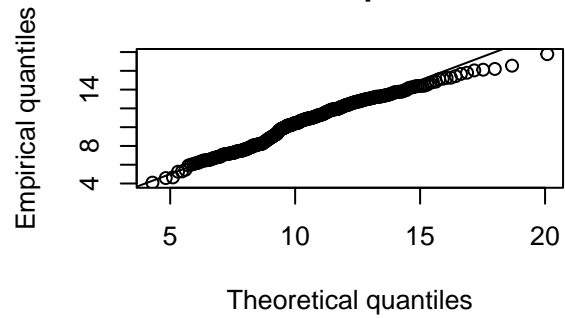
```
## Fitting of the distribution ' norm ' by maximum likelihood
## Parameters :
##      estimate Std. Error
## mean 10.453201 0.14039999
## sd    2.569742 0.09927772
## Loglikelihood: -791.5192   AIC: 1587.038   BIC: 1594.667
## Correlation matrix:
##      mean sd
## mean  1  0
## sd    0  1
```

```
plot(gamma_)
```

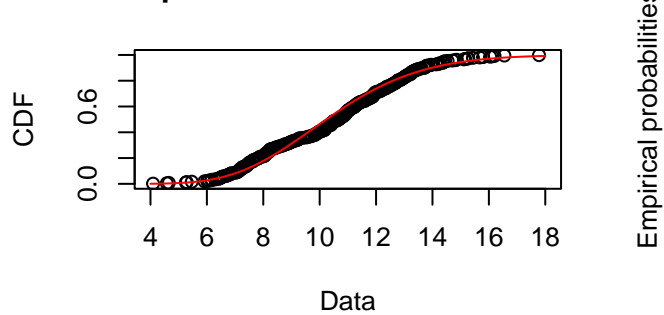
**Empirical and theoretical dens.**



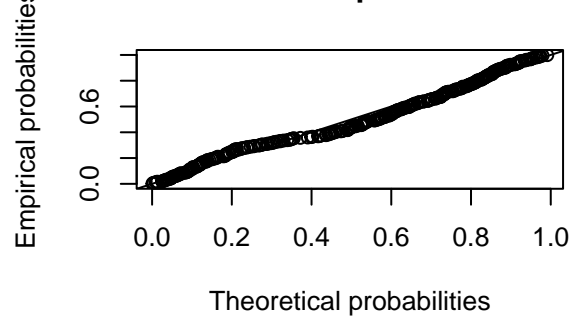
**Q-Q plot**



**Empirical and theoretical CDFs**






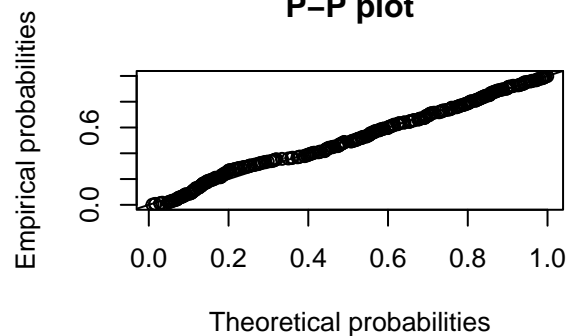
**P-P plot**



```
summary(gamma_)
```

```
## Fitting of the distribution ' gamma ' by maximum likelihood
## Parameters :
##      estimate Std. Error
## shape 15.54936  1.1887735
## rate   1.48756  0.1155788
## Loglikelihood: -794.6766   AIC:  1593.353   BIC:  1600.981
## Correlation matrix:
##      shape      rate
## shape 1.0000000  0.9839713
## rate   0.9839713  1.0000000
```

```
plot(weibull_)
```

**Empirical and theoretical dens.****Q-Q plot****Empirical and theoretical CDFs****P-P plot**

```
summary(weibull_)
```

```
## Fitting of the distribution ' weibull ' by maximum likelihood
## Parameters :
##      estimate Std. Error
## shape  4.554227  0.1927003
## scale 11.451692  0.1450187
## Loglikelihood: -791.1953   AIC:  1586.391   BIC:  1594.019
## Correlation matrix:
##      shape  scale
## shape 1.00000 0.32005
## scale 0.32005 1.00000
```

Is alpha significant for each hypothesis?

```
Field$Sex <- as.factor(Field$Sex)
```

```
parasiteLoad::getParamBounds("normal", data = Field, response = "WL")
```

```
##      L1start      L1LB      L1UB      L2start      L2LB      L2UB
## 10.453201105  4.091833908 17.777929279 10.453201105  4.091833908 17.777929279
##  alphaStart      alphaLB      alphaUB      mysdStart      mysdLB      mysdUB
##  0.000000000 -5.000000000  5.000000000  1.000000000  0.000000001 10.000000000
```

```
speparam <- c(L1start = 10,
              L1LB = 1e-9,
              L1UB = 20,
              L2start = 10,
```

```

L2LB = 1e-9,
L2UB = 20,
alphaStart = 0, alphaLB = -5, alphaUB = 5,
myshapeStart = 1, myshapeLB = 1e-9, myshapeUB = 5)

##All
fitWL_Sex <- parasiteLoad::analyse(data = Field,
                                   response = "WL",
                                   model = "normal",
                                   group = "Sex")

```

```

## [1] "Analysing data for response: WL"
## [1] "Fit for the response: WL"
## [1] "Fitting for all"
## [1] "Fitting model basic without alpha"
## [1] "Did converge"
## [1] "Fitting model basic with alpha"
## [1] "Did converge"
## [1] "Fitting model advanced without alpha"
## [1] "Did converge"
## [1] "Fitting model advanced with alpha"
## [1] "Did converge"
## [1] "Fitting for groupA : F"
## [1] "Fitting model basic without alpha"
## [1] "Did converge"
## [1] "Fitting model basic with alpha"
## [1] "Did converge"
## [1] "Fitting model advanced without alpha"
## [1] "Did converge"
## [1] "Fitting model advanced with alpha"
## [1] "Did converge"
## [1] "Fitting for groupB : M"
## [1] "Fitting model basic without alpha"
## [1] "Did converge"
## [1] "Fitting model basic with alpha"
## [1] "Did converge"
## [1] "Fitting model advanced without alpha"
## [1] "Did converge"
## [1] "Fitting model advanced with alpha"
## [1] "Did converge"
## [1] "Testing H0 no alpha vs alpha"
##      dLL dDF      pvalue
## 1 3.78   1 0.005963574
## [1] "Testing H1 no alpha vs alpha"
##      dLL dDF      pvalue
## 1 2.95   1 0.0151356
## [1] "Testing H2 groupA no alpha vs alpha"
##      dLL dDF      pvalue
## 1 1.37   1 0.09795211
## [1] "Testing H2 groupB no alpha vs alpha"
##      dLL dDF      pvalue
## 1 2.56   1 0.02360156
## [1] "Testing H3 groupA no alpha vs alpha"
##      dLL dDF      pvalue

```

```
## 1 1.61    1 0.07241473
## [1] "Testing H3 groupB no alpha vs alpha"
##      dLL dDF      pvalue
## 1 2.48    1 0.02590116
## [1] "Testing H1 vs H0"
##      dLL dDF      pvalue
## 1 0.95    1 0.168589
## [1] "Testing H2 vs H0"
##      dLL dDF      pvalue
## 1 0.19    3 0.9426523
## [1] "Testing H3 vs H1"
##      dLL dDF      pvalue
## 1 3.28    4 0.1608561
## [1] "Testing H3 vs H2"
##      dLL dDF      pvalue
## 1 4.04    2 0.01768312
```

```
parasiteLoad::analyse(data = Field,
                      response = "WL",
                      model = "normal",
                      group = "Sex")
```

```
## [1] "Analysing data for response: WL"
## [1] "Fit for the response: WL"
## [1] "Fitting for all"
## [1] "Fitting model basic without alpha"
## [1] "Did converge"
## [1] "Fitting model basic with alpha"
## [1] "Did converge"
## [1] "Fitting model advanced without alpha"
## [1] "Did converge"
## [1] "Fitting model advanced with alpha"
## [1] "Did converge"
## [1] "Fitting for groupA : F"
## [1] "Fitting model basic without alpha"
## [1] "Did converge"
## [1] "Fitting model basic with alpha"
## [1] "Did converge"
## [1] "Fitting model advanced without alpha"
## [1] "Did converge"
## [1] "Fitting model advanced with alpha"
## [1] "Did converge"
## [1] "Fitting for groupB : M"
## [1] "Fitting model basic without alpha"
## [1] "Did converge"
## [1] "Fitting model basic with alpha"
## [1] "Did converge"
## [1] "Fitting model advanced without alpha"
## [1] "Did converge"
## [1] "Fitting model advanced with alpha"
## [1] "Did converge"
## [1] "Testing H0 no alpha vs alpha"
##      dLL dDF      pvalue
## 1 3.78    1 0.005963574
## [1] "Testing H1 no alpha vs alpha"
```



```

##      dLL dDF      pvalue
## 1 2.95      1 0.0151356
## [1] "Testing H2 groupA no alpha vs alpha"
##      dLL dDF      pvalue
## 1 1.37      1 0.09795211
## [1] "Testing H2 groupB no alpha vs alpha"
##      dLL dDF      pvalue
## 1 2.56      1 0.02360156
## [1] "Testing H3 groupA no alpha vs alpha"
##      dLL dDF      pvalue
## 1 1.61      1 0.07241473
## [1] "Testing H3 groupB no alpha vs alpha"
##      dLL dDF      pvalue
## 1 2.48      1 0.02590116
## [1] "Testing H1 vs H0"
##      dLL dDF      pvalue
## 1 0.95      1 0.168589
## [1] "Testing H2 vs H0"
##      dLL dDF      pvalue
## 1 0.19      3 0.9426523
## [1] "Testing H3 vs H1"
##      dLL dDF      pvalue
## 1 3.28      4 0.1608561
## [1] "Testing H3 vs H2"
##      dLL dDF      pvalue
## 1 4.04      2 0.01768312

## $H0
##
## Call:
## bbmle::mle2(minuslogl = response ~ dnorm(mean = MeanLoad(L1,
##      L1, alpha, HI), sd = mysd), start = start, method = config$method,
##      optimizer = config$optimizer, data = data, lower = c(L1 = paramBounds[["L1LB"]],
##      mysd = paramBounds[["mysdLB"]], alpha = paramBounds[["alphaLB"]]),
##      upper = c(L1 = paramBounds[["L1UB"]], mysd = paramBounds[["mysdUB"]],
##      alpha = paramBounds[["alphaUB"]]), control = config$control)
##
## Coefficients:
##      L1      mysd      alpha
## 9.814002 2.540905 -0.281860
##
## Log-likelihood: -787.74
## Best method: bobyqa
##
## $H1
##
## Call:
## bbmle::mle2(minuslogl = response ~ dnorm(mean = MeanLoad(L1,
##      L2, alpha, HI), sd = mysd), start = start, method = config$method,
##      optimizer = config$optimizer, data = data, lower = c(L1 = paramBounds[["L1LB"]],
##      mysd = paramBounds[["mysdLB"]], L2 = paramBounds[["L2LB"]],
##      alpha = paramBounds[["alphaLB"]]), upper = c(L1 = paramBounds[["L1UB"]],
##      mysd = paramBounds[["mysdUB"]], L2 = paramBounds[["L2UB"]],
##      alpha = paramBounds[["alphaUB"]]), control = config$control)

```

```

##
## Coefficients:
##      L1      L2      alpha      mysd
## 9.5753580 10.0978117 -0.2526079 2.5337254
##
## Log-likelihood: -786.79
## Best method: bobyqa
##
## $H2
## $H2$groupA
##
## Call:
## bbmle::mle2(minuslogl = response ~ dnorm(mean = MeanLoad(L1,
##      L1, alpha, HI), sd = mysd), start = start, method = config$method,
##      optimizer = config$optimizer, data = data, lower = c(L1 = paramBounds[["L1LB"]],
##      mysd = paramBounds[["mysdLB"]], alpha = paramBounds[["alphaLB"]]),
##      upper = c(L1 = paramBounds[["L1UB"]], mysd = paramBounds[["mysdUB"]],
##      alpha = paramBounds[["alphaUB"]]), control = config$control)
##
## Coefficients:
##      L1      mysd      alpha
## 9.9286873 2.5773333 -0.2326945
##
## Log-likelihood: -395.07
## Best method: bobyqa
##
## $H2$groupB
##
## Call:
## bbmle::mle2(minuslogl = response ~ dnorm(mean = MeanLoad(L1,
##      L1, alpha, HI), sd = mysd), start = start, method = config$method,
##      optimizer = config$optimizer, data = data, lower = c(L1 = paramBounds[["L1LB"]],
##      mysd = paramBounds[["mysdLB"]], alpha = paramBounds[["alphaLB"]]),
##      upper = c(L1 = paramBounds[["L1UB"]], mysd = paramBounds[["mysdUB"]],
##      alpha = paramBounds[["alphaUB"]]), control = config$control)
##
## Coefficients:
##      L1      mysd      alpha
## 9.6835336 2.5023082 -0.3393935
##
## Log-likelihood: -392.47
## Best method: bobyqa
##
## $H3
## $H3$groupA
##
## Call:
## bbmle::mle2(minuslogl = response ~ dnorm(mean = MeanLoad(L1,
##      L2, alpha, HI), sd = mysd), start = start, method = config$method,
##      optimizer = config$optimizer, data = data, lower = c(L1 = paramBounds[["L1LB"]],
##      mysd = paramBounds[["mysdLB"]], L2 = paramBounds[["L2LB"]],
##      alpha = paramBounds[["alphaLB"]]), upper = c(L1 = paramBounds[["L1UB"]],
##      mysd = paramBounds[["mysdUB"]], L2 = paramBounds[["L2UB"]],

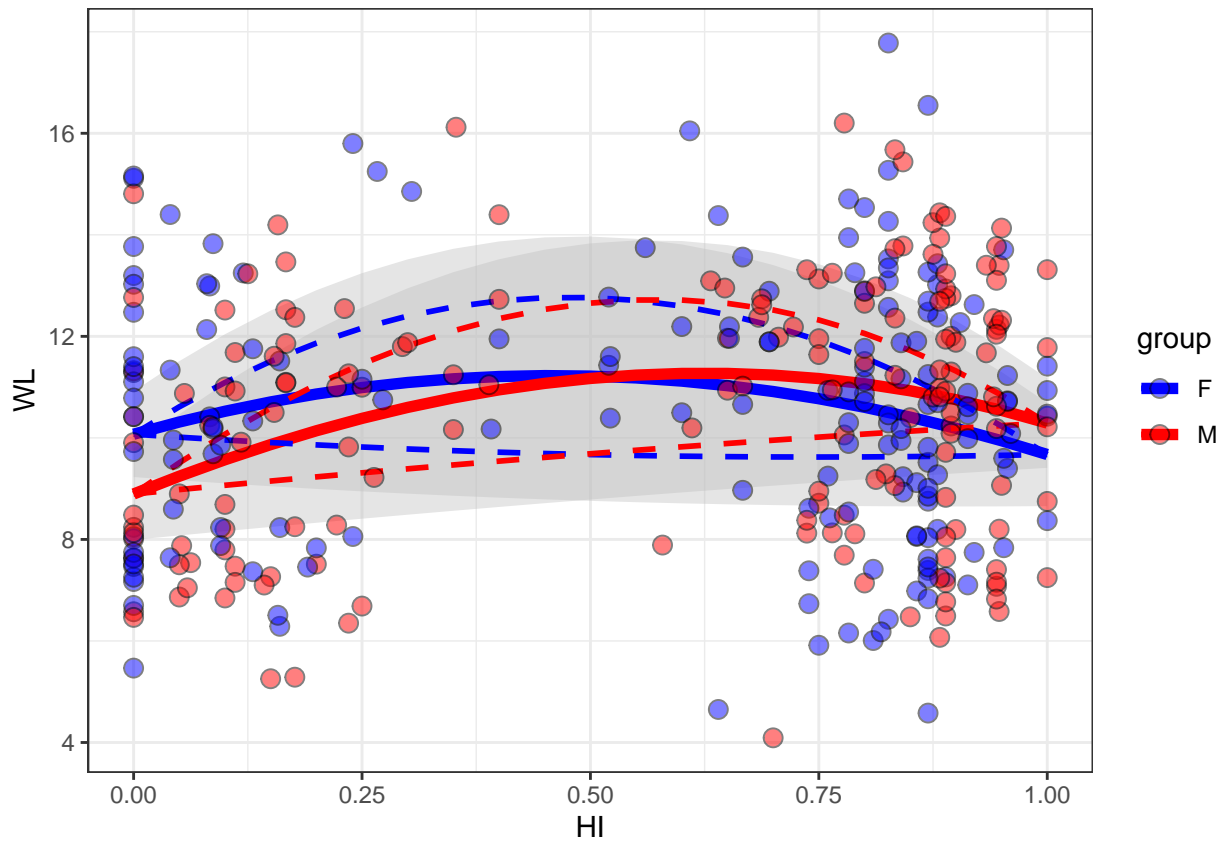
```

```

##      alpha = paramBounds[["alphaUB"]]), control = config$control)
##
## Coefficients:
##      L1      L2      alpha      mysd
## 10.0727199  9.6730112 -0.2723334  2.5735464
##
## Log-likelihood: -394.83
## Best method: bobyqa
##
## $H3$groupB
##
## Call:
## bbmle::mle2(minuslogl = response ~ dnorm(mean = MeanLoad(L1,
##      L2, alpha, HI), sd = mysd), start = start, method = config$method,
##      optimizer = config$optimizer, data = data, lower = c(L1 = paramBounds[["L1LB"]],
##      mysd = paramBounds[["mysdLB"]], L2 = paramBounds[["L2LB"]],
##      alpha = paramBounds[["alphaLB"]]), upper = c(L1 = paramBounds[["L1UB"]],
##      mysd = paramBounds[["mysdUB"]], L2 = paramBounds[["L2UB"]],
##      alpha = paramBounds[["alphaUB"]]), control = config$control)
##
## Coefficients:
##      L1      L2      alpha      mysd
##  8.8949407 10.2920779 -0.3288149  2.4464957
##
## Log-likelihood: -388.68
## Best method: bobyqa
plot_WL_Sex<- bananaPlot(mod = fitWL_Sex$H3,
      data = Field,
      response = "WL",
      group = "Sex") +
      scale_fill_manual(values = c("blue", "red")) +
      scale_color_manual(values = c("blue", "red")) +
      theme_bw()

## Scale for fill is already present.
## Adding another scale for fill, which will replace the existing scale.
## Scale for colour is already present.
## Adding another scale for colour, which will replace the existing scale.
plot_WL_Sex

```



H0: the expected load for the subspecies and between 2 groups is the same

H1: the mean load across 2 groups is the same, but can differ across subspecies

H2: the mean load across subspecies is the same, but can differ between the 2 groups

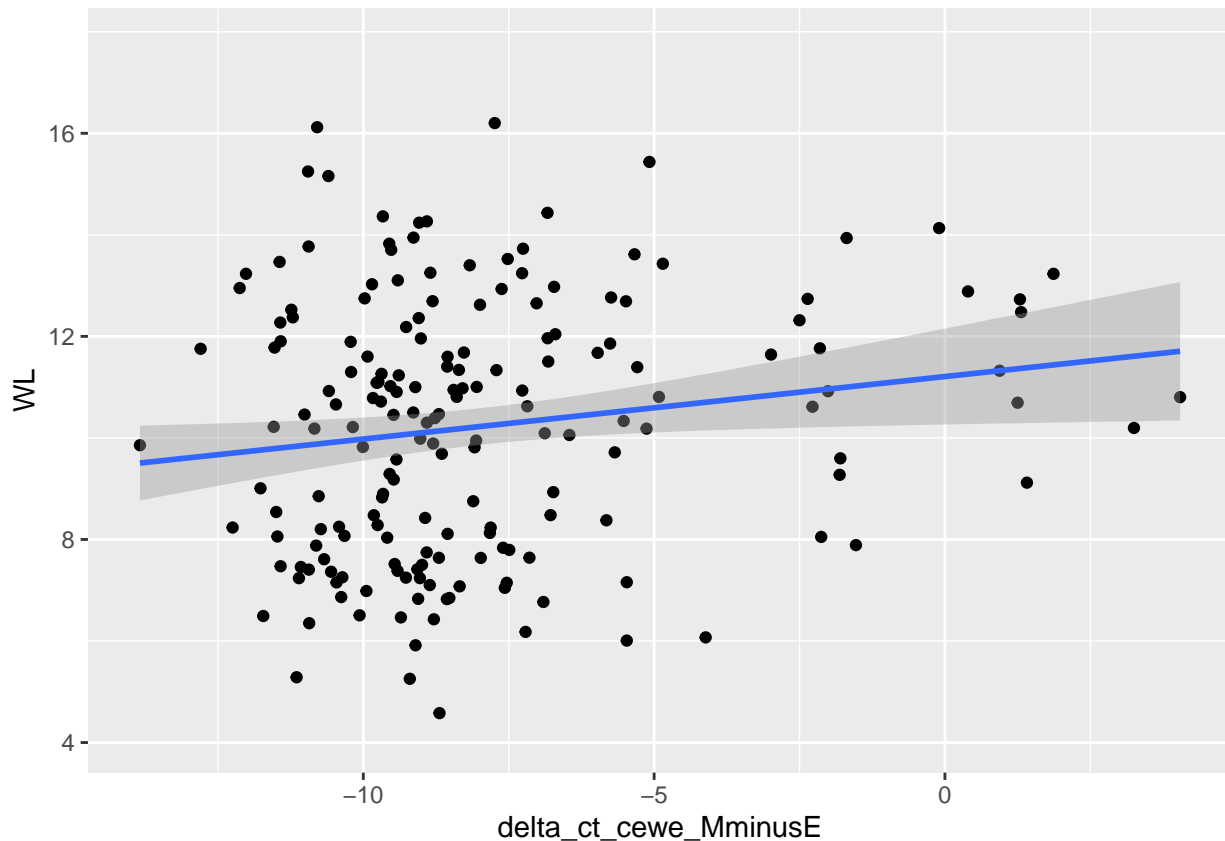
H3: the mean load can differ both across subspecies and between 2 groups

```
ggplot(data = Field, aes(x = delta_ct_cewe_MminusE, y = WL)) +
  geom_point() +
  stat_smooth(method= "lm")
```

```
## `geom_smooth()` using formula = 'y ~ x'
```

```
## Warning: Removed 150 rows containing non-finite values (`stat_smooth()`).
```

```
## Warning: Removed 150 rows containing missing values (`geom_point()`).
```



```
Field2 <- Field %>%
  drop_na(delta_ct_cewe_MminusE)

cor(Field2$WL, Field2$delta_ct_cewe_MminusE)

## [1] 0.1617751

tolerance <- lm(WL ~ delta_ct_cewe_MminusE, data = Field)

summary(tolerance)

##
## Call:
## lm(formula = WL ~ delta_ct_cewe_MminusE, data = Field)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -5.5596 -2.1168  0.2586  1.7938  6.2395
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    11.20837    0.47860   23.419  <2e-16 ***
## delta_ct_cewe_MminusE  0.12310    0.05551    2.218  0.0278 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.486 on 183 degrees of freedom
```


```
## (150 observations deleted due to missingness)
## Multiple R-squared: 0.02617, Adjusted R-squared: 0.02085
## F-statistic: 4.918 on 1 and 183 DF, p-value: 0.02781
```

```
confint(tolerance)
```

```
##                2.5 %    97.5 %
## (Intercept)    10.26408377 12.1526478
## delta_ct_cewe_MminusE 0.01358008 0.2326143
```

```
ggplot(data = Field, aes(x = OPG, y = WL)) +
  geom_point() +
  stat_smooth(method= "lm") +
  scale_x_log10()
```

```
## Warning: Transformation introduced infinite values in continuous x-axis
## Transformation introduced infinite values in continuous x-axis
## `geom_smooth()` using formula = 'y ~ x'
## Warning: Removed 280 rows containing non-finite values (`stat_smooth()`).
## Warning: Removed 157 rows containing missing values (`geom_point()`).
```



```
Field2 <- Field %>%
  drop_na(OPG)

cor(Field2$WL, Field2$OPG)
```

```
## [1] 0.04972871
```

```
tolerance <- lm(WL ~ OPG, data = Field)
```

```
summary(tolerance)
```

```
##
## Call:
## lm(formula = WL ~ OPG, data = Field)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -6.4814 -2.3628  0.1324  1.9028  7.1916
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 1.057e+01  1.986e-01  53.247  <2e-16 ***
## OPG         2.278e-08  3.449e-08   0.661    0.51
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.616 on 176 degrees of freedom
## (157 observations deleted due to missingness)
## Multiple R-squared:  0.002473,    Adjusted R-squared:  -0.003195
## F-statistic: 0.4363 on 1 and 176 DF,  p-value: 0.5098
```

```
confint(tolerance)
```

```
##              2.5 %      97.5 %
## (Intercept) 1.018133e+01 1.096510e+01
## OPG         -4.527973e-08 9.083863e-08
```